

Tensor Decomposition and Memory Networks for SAT Reading Comprehension

Saahil Madge

2016

Advised by Professor Christiane Fellbaum

SUBMITTED TO PRINCETON UNIVERSITY
DEPARTMENT OF COMPUTER SCIENCE

This Thesis represents my own work in accordance with University regulations

Date of Submission: April 29, 2016

Abstract

We present a general approach for machine comprehension tasks by converting the text to a knowledge graph and the questions to queries on the graph. We extend [19] and use the Stanford NLP Toolkit’s Dependency Parser [17, 9] to transform each sentence into a set of entity-relation triples. We use word2vec [18] to convert the questions into queries on the graph. We present a tensor decomposition approach to answering queries by adding Semantically Smooth Embedding [11] to RESCAL [20]. We also generalize the Memory Networks [28, 25] architecture to take any knowledge graph as input. We evaluate these models on three full SAT reading comprehension tests. The models presented here outperform their respective baselines. Both models demonstrate the ability to capture the semantic and structural information in the text and answer questions using that information.

Acknowledgements

I would like to say thank you to Professor Christiane Fellbaum for her guidance and encouragement over the past year in helping a student who didn't know anything about linguistics teach a computer to read and answer SAT questions. Professor Fellbaum, your willingness to meet with me every week, teach me concepts, and work through problems with me is the main reason I could take this thesis from a mere idea to its finished state. Your advice about work, graduate school, and life in general has been very helpful too. A big thank you to Professor Emmanuel Abbe as well for helping me understand knowledge graphs and for agreeing to review this thesis.

Of course a huge thank you to all my friends here who have helped me learn and grow over the past four years. Thank you to all the members of The Lodge, who made sure that not a single day at Princeton was a boring one, and with whom I'm lucky to have spent so much time. We've lived together for four years and we'll be best friends for a lot longer. Thanks to everyone at Charter for all the great dinner conversations, gaming sessions, and wonderful memories. Thanks to the HackPrinceton team for helping me out with the hardest of tasks. Thanks to all my other friends who haven't been mentioned yet but whom I value just as much.

Most importantly, a tremendous, tremendous thanks to my mother, father, and sister. It is impossible to say in words how grateful I am for everything you've done for me. All I have ever accomplished has been thanks to your love and support every step of the way. I love you all so much.

Contents

1	Introduction	6
2	Technical Review	9
2.1	Representation Techniques	9
2.1.1	Bag-of-words	9
2.1.2	Dependency Parsing	10
2.1.3	word2vec	11
2.1.4	Knowledge Graph Overview	11
2.2	Neural Networks	12
2.2.1	Feed-forward Neural Networks	13
2.2.2	Recurrent Neural Networks	14
3	Related Work	15
4	Our Model	17
4.1	Knowledge Graph Representation	18
4.1.1	Advantages of Knowledge Graph Representation . .	18
4.1.2	Converting to Knowledge Graph Representation . .	21
4.2	Tensor Decomposition	22
4.2.1	Knowledge Graph Tensor	22
4.2.2	Tensor Factorization Overview	23
4.2.3	Semantically Smooth Embedding	24
4.2.4	Solving the Optimization Problem	26
4.2.5	Answering Queries	28
4.3	Memory Networks	29
4.3.1	Memory Network Architecture	29
4.3.2	Memory Network Enhancements	31
4.4	Model Comparison	32
5	Data	33
6	Evaluation and Discussion	35
6.1	Implementation Details	35
6.2	Evaluation Process	35
6.3	Results	36

6.4	Comparison to Baseline	37
6.5	Interpretation of Results	38
6.6	Performance by Difficulty	40
6.7	Comparing Tensor Model to Memory Networks	43
7	Future Work	44
7.1	Linguistic Analysis	44
7.1.1	Knowledge Graph Extraction Improvements	44
7.1.2	Question Parsing	45
7.1.3	Choosing an Answer	46
7.2	Tensor Model Improvements	46
7.3	Memory Network Improvements	47
8	Conclusion	47
9	Appendix	51
9.1	MCTest Example	51
9.2	SAT Example	52
9.3	Derivation of Output Layer Error	53
9.4	Full Evaluation Data	55

1 Introduction

Machine question-answering (QA) has been a popular subject in recent research. It is a very interesting field from an academic as well as practical perspective, requiring knowledge and techniques across the areas of artificial intelligence, machine learning, natural language processing, linguistics, and mathematics.

QA has applications to many aspects of daily life. A famous use of question-answering techniques is in Apple's Siri program. Siri combines speech recognition software with question-answering techniques to create a virtual iPhone assistant. Another example is IBM's Watson, which answered questions in many subjects and performed quite well as a contestant on the game show *Jeopardy!* A less spectacular but more commonly-used application can be seen by typing in a simple question such as *How many calories are in an apple?* on Google. The search results provide a detailed nutritional breakdown of an apple.

The general term "QA" covers a broad array of subfields. Information retrieval, factual question-answering, reading comprehension, and inference are all part of the greater question-answering area. QA is a rapidly evolving field, but as the examples above indicate, recent research has focused extensively on answering spoken questions and factual questions.

However, the subfield of machine comprehension has not been studied nearly as well, though it has become quite popular in the past few years. Machine comprehension involves training programs and systems to read text and answer questions based on the newly acquired information. In most other QA tasks, the system has a pre-stored large database of facts (called a *knowledge base*). It first interprets the question and uses the database to provide the relevant factual answer. Machine comprehension tasks, on the other hand, require the system to parse both the informational text as well as the question, and then provide the correct answer. As there is a possibility of error within this dynamically created knowledge base, machine comprehension tasks are some of the most difficult QA tasks.

Machine comprehension has a wide range of applications. Reading comprehension tests are a natural choice, but fields like document retrieval also use machine comprehension techniques. Financial firms that use news analysis to make trading decisions heavily rely on these techniques as well. Many long-term goals of AI such as dialogue and humanoid robots cannot be achieved without significant advance in this area [27]. Another incentive to study machine comprehension is that new research here can drive research in many other areas of artificial intelligence and natural language processing.

As a consequence of machine comprehension tasks being so difficult, most research with reading comprehension tests has focused on elementary school level tests. The most common dataset is MCTest [23], which we discuss in more detail in Section 3. This dataset is a set of passages and reading comprehension questions created by crowdsourcing. The stories are “carefully limited to those a young child would understand” [23]. An example of an MCTest passage and question can be found in Appendix Section 9.1. Another common dataset is Facebook’s bAbi dataset [27], which has a collection of short toy problems that can be used to train machine comprehension systems. Berant et al. [3] train a system to answer questions based on a paragraph describing biological processes, which is certainly an advanced topic, but the system is specialized for biological tasks and not applicable to general reading comprehension.

In this paper we focus on SAT reading comprehension questions. The major reason is that they are far more complex than MCTest questions and other commonly-used machine comprehension datasets. The sentence structures, vocabulary, and topics covered in SAT passages are all far more varied than in the other benchmarks. The questions are much harder as well, and understanding the question and the individual choices is already a difficult task. SAT questions often test underlying themes and broad ideas rather than particular factual details (“Why” or “How” questions). In contrast, MCTest questions are often of a “Who”, “When”, or “What” form, which have traditionally been easier for programs to answer. An example of an SAT reading comprehension passage can be found in Appendix Section 9.2.

Another aspect that makes SAT questions so difficult is that they require inference to answer, as opposed to simply matching words or syntax. MCTest and similar questions can mostly be answered through word-matching (discussed in Section 2.1.1) or via syntactic matching. Word-matching means the system will try to find which sentences in the passage have the most words in common with the query, and use that to choose an answer. Syntactic matching is the same principle, but aims to match syntactic structure. SAT questions, however, almost never repeat the answer words in the question itself, and the query structure is not correlated with passage structure. To solve these questions, our system has to actually understand the high-level concepts and information that is implied but not directly on the surface. This is a significantly harder problem than what previous research has tackled, but solving it is necessary to continue moving forward in the field of machine comprehension. By trying to answer SAT questions we can attempt to solve problems that will be present in real-world applications but for which there are no existing approaches.

Due to the difficulty of interpreting the text as-is, we represent the text using relational logic instead. We dynamically create a *Knowledge Graph* (covered in more detail in Section 2.1) as we parse through the text, and frame each question as a query on this knowledge graph. The knowledge graph representation is typically used for factual data that has already been collected in entity-relation triple form. As far as we know, no previous research has tried to dynamically create a knowledge graph for machine comprehension tasks.

Traditionally most research on relational learning has used either neural networks or tensor decomposition to answer queries on the graph. In this work we present an improved version of both methods. For tensor factorization we enhance a method known as RESCAL [1, 20] with a technique known as *Semantically Smooth Embedding* [11]. RESCAL trains a factorization of the knowledge graph using Alternating Least Squares, and this factorization can then be used to answer queries.

The neural network approaches such as TransE [6] and TransH [26] train an embedding method for the labels on the knowledge graph, and use the trained model to answer queries.

However, none of these previous approaches are designed for general question-answering. We enhance the Memory Networks architecture proposed by Weston et al. [28] and improved by Sukhbaatar et al. [25] to take knowledge graph triples and queries as input, and provide an answer as output. The existing models are discussed in Section 3 and our architecture is detailed in Section 4.3.

All models and data are available at <https://github.com/SaahilMadge/Senior-Thesis>. Our goal with this research is to present a way to combine ideas that were previously used in isolation for QA or other learning tasks into a system which can parse and interpret a complex text and then answer difficult high-level questions about the text. The passages and questions used here are more difficult than those in previous research, and the goals more ambitious, but we hope that this paper will encourage others to tackle problems of equal or greater ambition and drive progress in machine comprehension and related fields.

2 Technical Review

In this section we review some of the techniques used in previous research and in this paper.

2.1 Representation Techniques

We begin by reviewing various methods for embedding text and performing queries on the embedded representations.

2.1.1 Bag-of-words

Bag-of-words is a technique in which we treat each sentence as simply a collection of words, without paying attention to their order or the dependencies between individual words. If we have some vocabulary of size V , we can encode each sentence as a vector of size V where each vector element corresponds to a word in the vocabulary. If the word appears in the




Figure 1: A small example knowledge graph, from [9]

sentence then we store a 1 in the corresponding vector element, and if it does not appear we store a 0. If every word gets its own vector (so there only one element has a 1 and all others have 0) it is called *one-hot* vector encoding. Most designs encode each input sentence and the query as a sum of one-hot vectors. These vectors are not used as-is, but are first embedded into some dimension d . The learning algorithm is run on the embedded vectors, and we can output a response in a specified format. Two common methods are to return a probability vector of size V where each element represents how likely that word is to be the output, or to pick the most likely word from that distribution and return only that word.

2.1.2 Dependency Parsing

Dependency parsing is an important technique used to understand the structural information in a sentence. Dependency parsing creates a graph of the words in the sentence and describes how each word relates to other words, using the standard Universal Dependencies [10]. A sample graph is shown in Figure 1, from [9]. Each dependency in the graph has a *governor* and *dependent*. The collection of dependencies contains all the structural relationships in the sentence. Each dependency is labeled to show the exact relationship between governor and dependent. For example, in the graph in Figure 1 the label *nsubj* highlights the nominal subject of a particular clause in the sentence [10]. The specific parser we use is the Stanford High-Performance Dependency Parser by Chen and Manning [9].

Once the dependency graph is created, we can analyze it to extract the relevant information in the sentence. The exact details are discussed more in Section 4.1.2. The basic approach

is to find the main subject of the sentence and either the direct object it acts on or some description of the subject. This information can be turned into an entity-relation triple on the knowledge graph.

2.1.3 word2vec

The *word2vec* embedding technique introduced by Mikolov et al. [18] is perhaps the most common embedding technique used in NLP research. It is actually an umbrella term for various methods, such as *skip-gram* or even bag-of-words embedding. The word2vec technique produces a vector representation of words, often with several hundred elements. These representations can then be manipulated as vectors in the higher-dimensional space, allowing for interesting operations on words. For example, $\text{vector}(\text{"king"}) - \text{vector}(\text{"man"}) + \text{vector}(\text{"woman"})$ produces a vector that is very close to $\text{vector}(\text{"queen"})$. The reason word2vec is able to embed so successfully is generally thought to be the fact that the process of embedding gives similar words a similar embedding (measured by cosine similarity). Cosine similarity is also used to find the similarity between two arbitrary words. *word2vec.similarity*(word1, word2) is a function within word2vec which returns a measure between 0 (not at all similar) and 1 (identical).

2.1.4 Knowledge Graph Overview

A knowledge graph generally refers to a set of *entity-relation triples*. These triples take the form $\langle e_1, r, e_2 \rangle$ where e_1 and e_2 are entities and r describes the relation between these two entities. Entities and relations (collectively referred to as *labels*) can be repeated across triples. The total set of triples denotes all the knowledge in our “Knowledge Base”.

We can represent these triples as a directed graph, with each entity as a vertex and the relations as edge types between them. Figure 2 shows an example of a very small knowledge graph. Note that there is only one node per entity. This allows the graph to store multiple pieces of information regarding the same entity together, to allow for easy processing by

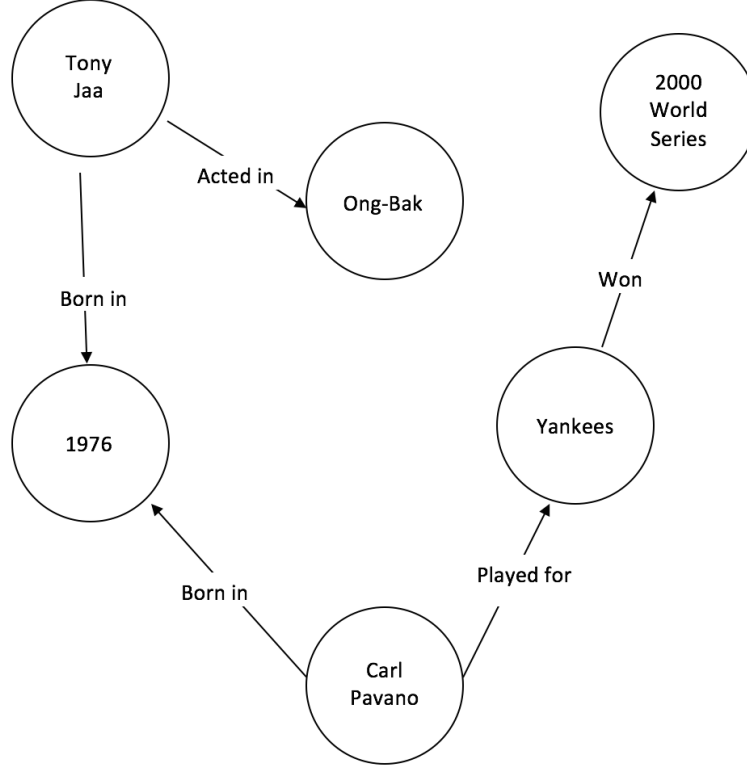


Figure 2: A small example knowledge graph

the program while also making it simple for humans to understand. The details of our knowledge graph model can be found in Section 4.1.

Once a knowledge graph is built, we can pose queries about the graph. For example, we can provide two entities e_1, e_2 and ask which relation r links the two on the graph. We can also provide an entity e_1 and the relation r , and ask for which other entity e_2 does the triple $\langle e_1, r, e_2 \rangle$ exist on the graph. Put another way, the queries are either *link prediction* queries or *entity prediction* queries.

2.2 Neural Networks

As discussed in Section 1, recent research in machine comprehension has extensively used neural nets (NNs) and their variants. Here we provide a quick review of neural network design. These topics are covered in far more depth in [5, 21].

2.2.1 Feed-forward Neural Networks

The most basic model is known as a *feed-forward* neural net. These have several layers composed of units (nodes). There is an initial input layer, a final output layer, and *hidden layers* in the middle. A unit in layer l takes as input the output of nodes in the previous layer (or the actual input if this is the input layer). It multiplies these by some weight matrix W^l , transforms it using an activation function g , and sends its output to the next layer (or as the final output if this is the output layer). Common activation functions are the *tanh* and *sigmoid* functions. The sigmoid function is defined as $\sigma(x) = \frac{1}{1 + e^{-x}}$. For now we will assume that our activation function is sigmoid.

Mathematically, we define the input to the current layer as x^{l-1} , where x_{ij}^{l-1} means that it is the input from unit j in the previous layer to unit i in the current layer. W_{ij}^l is the weight on this connection, defined as 0 if node i does not depend on node j . The output of node j is $\sigma(W^l x_j^{l-1})$. Typically the quantity $W^l x_j^{l-1}$ is known as z_j . We can also define bias vectors b such that $z_j = W^l x_j^{l-1} + b_j^l$. To summarize, we can write the output of some layer l in vector notation as

$$o^l = \sigma(W^l x^{l-1} + b^l)$$

We can propagate this forward starting at the input layer, through the hidden layers, and finally at our output layer. We can then train our network using the backpropagation algorithm. We have some cost function C that is a measure of the error of our output. We find the error at the output layer, and then change the weights at each previous layer, working backwards to the input layer. The key to backpropagation relies on updating the weights based on their partial derivatives with relation to C . In-depth derivations are provided in [5, 21].

2.2.2 Recurrent Neural Networks

Feed-forward networks are good tools, but do not have any form of state (memory). Here we review recurrent neural networks (RNNs) which have a hidden state h_t . In-depth derivations are provided in [7].

At step t in a recurrent neural network, we have input x_t and a hidden state h_{t-1} . We also have weight matrices U, W, V . We can define the current state $h_t = \sigma(Ux_t + Wh_{t-1})$. The output $\hat{y}_t = \text{softmax}(Vh_t)$, where $\text{softmax}(z_i) = \frac{e^{z_i}}{\sum_j e^{z_j}}$. The softmax function output is a probability distribution over the vector.

Using the softmax function has several advantages. Because it is a smooth function, we can take its derivative, making it easier to calculate the backpropagation equations. Additionally, most RNNs and variants use a loss function known as cross-entropy loss. If the model output is \hat{y} and the true output is y , then cross-entropy loss is defined as $C = -\sum_i y_i \log(\hat{y}_i)$, which is the quantity $-y_i \log(\hat{y}_i)$ summed over all output nodes. When we pair a softmax output with cross-entropy loss, the error at the output node i is simply $\hat{y}_i - y_i$. This is very convenient because it can be calculated rapidly, so it allows for fast training of the RNN. Training is normally a bottleneck, and by combining a robust output function and robust loss function we get a very elegant output error. We have included a full derivation of this result in Appendix Section 9.3.

One drawback of RNNs is that they cannot learn dependencies across a large time gap. This is problematic when dealing with machine comprehension, as the information required to answer a particular question may be scattered across sentences, and in the case of SAT questions even across multiple paragraphs. This inability to learn long-range dependencies is known as the *Vanishing Gradient* problem, first discovered by Hochreiter [13]. During backpropagation a given time step depends on the gradients of all the time steps before it, and this decreases exponentially fast [7]. As a result, states at a much later time than the current state do not contribute to the update rules at the current state. Memory Networks are designed to avoid this problem, as discussed in Section 4.3.

3 Related Work

The first major research in machine comprehension was conducted by Hirschman, et al. [12]. Their system “Deep Read” takes a story and a series of questions as input, and proposes answers by using bag-of-words techniques with other methods such as stemming and pronoun resolution layered on top. On a collection of remedial reading questions for grades 3-6, Deep Read answered about 40% of the questions correctly.

Most recent research has used the MCTest dataset [23] released by Richardson, et al. at Microsoft research. There are a total of 500 passages in MC500, each designed for elementary school students. The stories are fictitious, which means the answers can only be found in the story itself. Richardson et al. also provide two baseline implementations. The first is a simple bag-of-words model with a sliding window. It scores about 51% accuracy on the MC500 questions. The second implementation adds a distance-based scoring metric and reaches 56% accuracy. Both implementations score significantly higher on questions where the answer is contained in just one sentence than on questions where the answer requires information across multiple sentences.

Narasimhan and Barzilay [19] also work on the MCTest tasks. Their main insight is to create a task-specific discourse parser to capture specific discourse relations in the MCTest passages, while the prevailing method had been to use generalized off-the-shelf discourse parsers. They create three models of increasing complexity. The first assumes each question can be answered using just one sentence, and estimates a joint probability distribution across the question, query, and answer. The second adds in the joint distribution across a second sentence as well, to handle the case in which the answer needs two sentences to answer. The third incorporates discourse relations between the two sentences to better understand the relationship between information in each sentence. The relations are defined as *Causal*, *Temporal*, *Explanation*, and *Other*. Most systems perform poorly on these types of questions. By modeling these explanatory relations the third model achieves almost 64% accuracy, easily outscoring the other two as well as the two MCTest baseline implementations.

Berant et al. [3] also focus on analyzing inter-sentence relations, with application specifically to passages and questions about biological processes. These passages generally describe a chemical reaction or other process in which there are various starting entities which interact with each other and form a new output. Understanding how the inputs interact and tracing the flow of the process is crucial to answering the question. To solve this, Berant et al. define events (or non-events) as *triggers*, and try to find relationships between events. The events can be thought of as nodes on a graph, with an edge defining some relation between the two. There are eight possible relations, including *cause*, *enable*, and *prevent*. They first identify events and then predict relations between the events. The queries are formulated as a graph and categorized as dependency questions, temporal questions, or true-false questions. This model scores almost 67% on the dataset, over 6% better than the next-best model.

Most recent machine comprehension research has focused on using neural nets as the core system. Neural nets are a very generalizable technique and in the past decade have become very popular. In fact, advances in neural net techniques have greatly contributed to the recent surge in machine comprehension research.

Weston et al. [28] introduced memory networks. These are a type of neural network which simulate a long-term memory. We discussed in Section 2.2.2 how RNNs have a “vanishing gradient” problem, and are not able to take advantage of states from more than a few steps prior. Memory networks have a memory module which stores old memories and then updates them given new inputs. When given a query, the memory network finds relevant memories, then finds memories that are relevant given those memories, and so on. Finally, it provides an answer to the query. Sukhbaatar et al. [25] improved this model by creating a model that can be trained end to end, while in the original design each module needed to be trained independently. This is the model that is the basis for our design, so we discuss it in more detail in Section 4.3. Kumar et al. [16] create “Dynamic Memory Networks”. Their design uses two memory modules. The semantic memory module stores

general knowledge while the episodic memory module iteratively finds memories relevant to the query.

Knowledge Graphs are also a popular research area. There are two main approaches that are used for solving knowledge graph embedding and query problems. The first is the neural network embedding style, popularized by Bordes et al. [6] and improved by Wang et al. [26]. Bordes’ TransE and Wang’s TransH both train a neural network to recognize triple embeddings in some higher dimension. The triples are embedded in such a way that triples that are “similar” according to some metric are embedded near each other. Queries can be answered more easily using the embedded vectors.

The second approach focuses on representing the knowledge graph as a tensor. Nickel et al. [20] proposed RESCAL, a model which converts the knowledge graph into a tensor of adjacency matrices and trains a latent rank- r factorization of the tensor using a technique called ASALSAN [1]. Queries about the knowledge graph can be easily converted into a search over a relevant subset of elements in the factorization. Chang et al. [8] created TRESICAL, which is an optimized version of RESCAL.

4 Our Model

This research presents three fundamental insights.

1. We dynamically convert the input text into a knowledge graph and transform the questions into queries on the graph. We answer the original question by finding an entity or relation that completes the query and matching it to the best answer choice.
2. We use tensor decomposition to embed the knowledge graph and questions. We make the RESCAL [20] more accurate by adding *Semantically Smooth Embedding*(SSE), as proposed by Guo et al. [11]. In their original paper they propose adding this type of embedding to tensor decomposition as future work, so to our knowledge we are the first to perform tensor decomposition with SSE.

3. We vectorize the knowledge graph and questions and pass them as input to a Memory Network. Given a set of embedded inputs and queries, this type of neural network takes advantage of “memory” to find new facts that are relevant given the current set of relevant facts, but may not have been relevant solely based on the query.

These techniques were first introduced in previous research. However, all of them were applied to separate domains. Neural nets have been used for machine comprehension on MCTest, bAbi, and questions of similar difficulty. Their embedding model is quite simple, however (often just bag-of-words). Knowledge graphs and tensor decomposition have traditionally been used only for relational learning when a knowledge graph or set of triples already exists. Combining these original models requires significant modification of their conceptual and mathematical basis.

4.1 Knowledge Graph Representation

The first part of our design is to dynamically convert the input passage into a knowledge graph and the questions into queries on the graph.

4.1.1 Advantages of Knowledge Graph Representation

There are several advantages to using a knowledge graph representation. The main benefit is that it actually preserves the meaning of the text. Whereas bag-of-words loses the semantic and structural information in a sentence, a knowledge graph representation almost completely stores the intended meaning of the original text. Most sentences in English can be broken down into Subject-Verb-Object triples, which require little if any processing to convert into the $\langle entity_1, relation, entity_2 \rangle$ format. The knowledge graph itself is simply a list of these triples. If we can parse the input text and extract the Subject-Verb-Object relationship, we can convert each sentence into a triple of the form

$$\langle subject, verb, object \rangle \Rightarrow \langle e_1, r, e_2 \rangle$$

The details of the conversion are covered in Section 4.1.2.

Even more importantly, the knowledge graph inherently has *memory*. What this means is that it automatically pieces together information about the same concept, even if that information occurs across different sentences or even different paragraphs. Consider a passage about a boy named Harry. Even if in one paragraph the text says that Harry was born in 1980 and several paragraphs later we learn that he has a friend named Ron, these relations both link to the single Harry entity node. By analyzing the outbound and inbound relations of the Harry entity node, we can easily find all the information the text has stated about Harry. Any question asked about Harry can be answered using this information. The knowledge graph representation is the best way for our system to easily find all information about an entity that was given in the original text.

Entities do not have to be concrete things; they can just as easily be abstract concepts. For example, Figure 2 shows entities like Carl Pavano (a person), but it also shows an entity node for the year 1976. To roughly generalize, any noun can be an entity node. The phrase “abstract concept” itself could be an entity node. Entity-relation triples can all be analyzed similarly, so abstract reasoning is as simple as concrete reasoning. Not only does a knowledge graph allow us to perform abstract reasoning, but it is arguably the only way to perform reasoning of this form.

A good way to understand the knowledge graph representation is to interpret it as a *concept map*. Every sentence in the text introduces a new concept (concrete or abstract) and/or provides more information about a previously introduced concept. All information is provided in such a way that the concept is either the actor or receiver of some action or description. We update our concept map with this new information by adding entity nodes as needed and the relation provided in the current sentence. Once the input text is parsed we have the full concept map, which gives complete information on each concept and the information we know about it.

Another benefit of using knowledge graphs is that we can use the existing theory and methods of analysis for these graphs. As mentioned in Section 1, passages as complex as SAT reading comprehension tests have rarely been analyzed semantically. Moreover, the fact that information comes in a variety of sentence structures across several paragraphs means that traditional semantic analysis will not suffice. Traditional analysis often stores information in the scope of its sentence or paragraph, which is not enough for our particular use case. SAT passages simply cannot be analyzed using these existing techniques. We must convert it to a form which we know how to analyze. Not only have knowledge graphs been studied extensively in the field of relational learning, but they can also be analyzed using techniques from graph theory. Essentially, we turn the problem of analyzing complex text into a problem of analyzing a knowledge graph.

At the same time, the knowledge graph representation also preserves human readability. A major problem with text input is that it is easily interpretable by humans but requires much preprocessing before it can be interpreted by a program. It is not necessary for our processed text to be in a format humans can interpret, but it certainly helps us reason about the problem and come up with interesting solutions. The entity-relation triple format is almost as intuitive to humans as the original textual input.

Of course, even the knowledge graph representation has some drawbacks. Because we build the knowledge graph using the statements in the original text, we can only perform analysis on the textual information. We can trace relations through the graph but there is no way to perform higher-level reasoning. For example, questions that ask about classification of new information or underlying themes of the passage cannot be answered. Nevertheless, these are flaws in any form of representation. The knowledge graph provides the most advantages and fewest drawbacks of any representation.

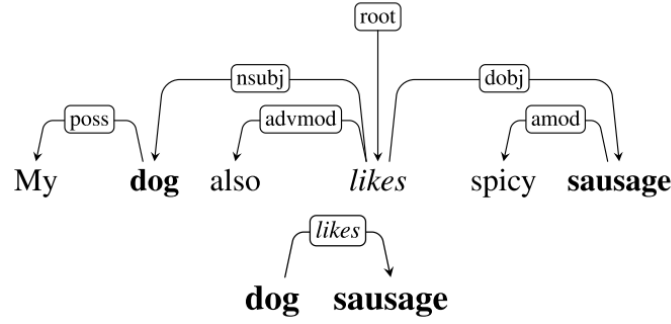


Figure 3: Extracting an entity-relation triple from a sentence, from [19]

4.1.2 Converting to Knowledge Graph Representation

In this section we present the details of converting the text into a Knowledge Graph. Section 5 discusses how the raw tests are converted into textual information that we can then parse.

We use the Stanford High Performance Dependency Parser [9] to extract the entity-relation information from the text. Section 2.1.2 explains how this tool creates a dependency graph of a particular sentence. We then provide a set of semantic rules to isolate the important information and create a triple from it. Figure 3 from [19] shows a simple sentence, as well as the extracted triple.

The information extracted should conform to knowledge graph format, namely $\langle entity_1, relation, entity_2 \rangle$. We want to frame the information in the sentence so that it is one entity with some relation to another entity. Most sentences are of the form where the subject performs some action or has some description, so we focus on these two types of sentences. In each sentence we isolate the main subject using the *nsubj* dependency. Then we check for one of two relations. If the subject is performing some action then the sentence can generally be extracted as $\langle subject, action, direct\ object \rangle$. In this case the *dobj* dependency is used to extract the action verb as well as the direct object. Figure 3 shows this case. In the second case, the sentence is often of the form “<subject> is... <description>”. Here we use the *cop* dependency to extract the verb and the description. The triple $\langle subject, copular\ verb, description \rangle$ is added to the knowledge graph. In English,

“verbs of being” such as “is”, “are”, “were”, etc. are all copular verbs. Narasimhan and Barzilay [19] do a similar type of extraction, but they create what they call an *entity graph*, which is slightly different from our knowledge graph. Their entity graph refers to linguistic entities from the dependency parsing, while our knowledge graph refers to information, not linguistic objects.

4.2 Tensor Decomposition

First we discuss the tensor decomposition approach to answering SAT questions. The advantage of using tensor decomposition is that it is based on well-understood mathematics, rather than simply empirical validation as with the neural network.

Our model is an improved version of RESCAL, proposed by Nickel et al. [20] and used by Chang et al. [8]. RESCAL itself takes the general tensor decomposition model proposed by Bader et al. [1] and applies it to relational learning.

4.2.1 Knowledge Graph Tensor

The tensor decomposition approach treats the knowledge graph as a directed graph. A tensor is simply a multi-dimensional array. In our case, we use a three-dimensional array \mathcal{X} to represent the knowledge graph. If we have N entities e_1, e_2, \dots, e_N and M relations k_1, k_2, \dots, k_M then we can consider \mathcal{X} as M layers of $N \times N$ matrices. Each layer \mathcal{X}_k is an adjacency matrix that marks which entities are linked by relation k . If we denote the set of triples as KB , then \mathcal{X} is formulated as

$$\mathcal{X}_{ijk} = \begin{cases} 1 & \text{if } \langle e_i, k, e_j \rangle \in KB \\ 0 & \text{otherwise} \end{cases}$$

Figure 4 from [8] presents a graphical view of \mathcal{X} .




Figure 4: A graphical representation of \mathcal{X} , from [8]

4.2.2 Tensor Factorization Overview

Our approach aims to factorize the tensor and take advantage of the inherent structure of the graph. To that end we train a rank- r factorization of each slice, as suggested in [1, 20].

$$X_k \approx AR_k A^T \text{ for } k \in [1, m]$$

Here, A is an $n \times r$ matrix which contains the latent-component representation of the entities in the knowledge graph. We can think of this as an r -dimensional embedding of the entities, where each entity has r hidden features. This is conceptually similar to the word2vec embedding, where each feature of a word vector does not necessarily mean something intuitive but is a hidden component found through training.

R_k is an $r \times r$ matrix which represents how these hidden features interact with each other for the k -th relation. R_k is asymmetric, which means that the hidden feature interactions are one-way. It is not necessarily true that $R_{ijk} = R_{jik}$, although it could be. Note that R itself is a three-dimensional tensor of shape $r \times r \times k$.

We decompose the tensor so that we can answer queries using the factorization, which has the hidden components of the graph and is thus more generalizable, as opposed to using the original tensor which is simply an adjacency tensor of the knowledge graph. For example, a

particular connection between two entities might not be explicit on the knowledge graph but can be uncovered by the factorization.

We frame the problem of finding A and R_k as an optimization problem, where we aim to minimize the difference between the original matrix \mathcal{X}_k and its factorization AR_kA^T . Specifically, define the loss function f as the squared Frobenius norm of this difference.

$$f(A, R_k) = \frac{1}{2} \sum_k \left\| \mathcal{X}_k - AR_kA^T \right\|_F^2 \quad (1)$$

where the Frobenius norm of some $a \times b$ matrix X is defined as

$$\|X\|_F = \sqrt{\sum_{i=1}^a \sum_{j=1}^b |X_{ij}|^2}$$

To find the optimal A and R_k we minimize Equation (1) along with some regularization function $g(A, R_k)$. Formally,

$$\min_{A, R_k} f(A, R_k) + \lambda g(A, R_k) \quad (2)$$

In [20][8] the regularization term is simply defined as $g(A, R_k) = \frac{1}{2} \lambda (\|A\|_F + \sum_k \|R_k\|_F)$. This is a standard regularization function where λ is a hyperparameter. The aim is to prevent overfitting when we try to solve the optimization problem. We enhance their approach by using a stronger regularization function which adds more constraints on the factorization.

4.2.3 Semantically Smooth Embedding

Semantically Smooth Embedding (SSE) was proposed by Guo et al. [11] as a way to improve the TransE [6] embedding model. Whereas the traditional models embed each individual fact in isolation, SSE imposes geometric constraints on the whole model.

Specifically, we modify the *Locally Linear Embedding* (LLE) constraint. This constraint states that each node should be roughly a linear combination of its neighbors [24]. Guo et al.

[11] created a version meant for neural net embeddings such as TransE but is not applicable to the tensor decomposition method. Here we review their version first and then present our new version for tensor decomposition.

In the original implementation each entity was labeled with a category and then embedded into a vector representation. All nodes in the same category as e_i were defined as its neighbors and this set was denoted $\mathcal{N}(e_i)$. Define

$$w_{ij} = \begin{cases} 1 & \text{if } e_j \in \mathcal{N}(e_i) \\ 0 & \text{otherwise} \end{cases}$$

$$R = \sum_{i=1}^n \left\| \mathbf{e}_i - \sum_{e_j \in \mathcal{N}(e_i)} w_{ij} \mathbf{e}_j \right\|_2^2$$

where R is the regularization term added to the loss function. The logic here is that we suffer more penalty in our loss function when entities that are close to each other in the original set (i.e. are neighbors) have embedded vectors that are far from each other. We cannot use this equation for our purposes since our entities do not have categories and we do not have embedded vectors, but we can apply similar logic.

We propose defining the neighbor set of each entity e_i in our graph as all other entities for which there is a directed edge from e_i to e_j in the knowledge graph. The full neighbor set matrix is denoted W .

$$w_{ij} = \begin{cases} 1, & \sum_k \mathcal{X}_{ijk} \geq 1 \\ 0, & \text{otherwise} \end{cases}$$

$$e_j \in \mathcal{N}(e_i) \text{ iff } w_{ij} = 1$$

Put another way, $w_{ij} = 1$ if and only if we have some relation r for which the triple $\langle e_i, r, e_j \rangle$ exists in the knowledge base. We do not consider entities which are related to e_i in such a way that e_i is the recipient of the relation. Note that this means the neighbor relation goes one way. $e_j \in \mathcal{N}(e_i)$ does not imply that $e_i \in \mathcal{N}(e_j)$.

Now consider the nodes in $\mathcal{N}(e_i)$. Since e_i is related to all of these, and because we want to enforce some geometric structure on our representation, we claim that the latent properties of each node e_i as represented in the matrix A should be some linear combination of the properties of its neighbor set.

$$A_i \approx \sum_{j=1}^n w_{ij} A_j = W_i A$$

We can penalize more for an A in which this distance is large. The reason we made the neighbor relation only one way is that we do not expect a large number of disconnected components of our knowledge graph, and a two-way neighbor set with this regularization would cause every node to converge to the same latent representation.

4.2.4 Solving the Optimization Problem

The new, semantically smooth regularization function is

$$g(A, R_k) = \sum_{i=1}^n \left\| A_i - \sum_{j=1}^n w_{ij} A_j \right\|_F^2 = \sum_{i=1}^n \|A_i - W_i A\|_F^2$$

We are now ready to actually solve for A and R_k . We want to minimize Equation (2). This is a nonlinear, non-convex optimization problem which requires complex techniques to solve. Calculating f and g is expensive, so we use the Alternating Least Squares(ALS) method [15]. Specifically, we use the ASALSAN factorization method [1].

ALS turns the non-convex optimization problem into a convex problem by fixing one of the unknowns. Then the problem can be solved with ordinary least-squares optimization, so we use the normal method of least-squares to update the unknown variable. Then we alternate by fixing this unknown and changing the other variable. The ordinary least-squares can be solved relatively fast, though calculating inverse matrices is expensive, and each unknown is calculated independently while keeping the other fixed, allowing the ALS algorithm to be easily parallelized. These two aspects make it very computationally fast, a quality necessary for our purposes.

We define

$$\bar{\mathcal{X}} = A\bar{R}(\mathbf{I}_{2m} \otimes A^T)$$

where \otimes is the Kronecker product and

$$\bar{\mathcal{X}} = (\mathcal{X}_1\mathcal{X}_1^T \dots \mathcal{X}_m\mathcal{X}_m^T)$$

$$\bar{R} = (R_1R_1^T \dots R_mR_m^T)$$

Then by ASALSAN we derive the following update rules. We first update A while keeping R_k fixed.

$$A \leftarrow \left[\sum_k \mathcal{X}_k A R_k^T + \mathcal{X}_k^T A R_k \right] \left[B_k + C_k \right]^{-1}$$

where

$$B_k = R_k A^T A R_k^T, \quad C_k = R_k^T A^T A R_k$$

To update R_k , we first fix A . Then we notice that if we vectorize \mathcal{X} and R_k , we can rewrite Equation (2) as

$$f(R_k) = \|\mathbf{vec}(\mathcal{X}_k) - (A \otimes A) \mathbf{vec}(R_k)\|$$

where we have removed an extra term $\lambda \|\mathbf{vec}(R_k)\|$ from the RESCAL model because R_k is no longer in our regularization term. This rewritten f is actually a linear regression problem whose solution is

$$R_k \leftarrow \left((A \otimes A)^T (A \otimes A) \right)^{-1} (A \otimes A) \mathbf{vec}(\mathcal{X}_k)$$

We alternate updating A and R_k until we have performed some maximum number of iterations or until $\frac{f(A, R_k)}{\|\mathcal{X}_k\|_F^2}$ is below some tolerance ϵ . The final factorization $\hat{\mathcal{X}}$ can now be used to answer queries.

4.2.5 Answering Queries

The first step in answering SAT questions with tensor decomposition is to convert them to queries on the knowledge graph. Queries are either link prediction queries or entity prediction queries (Section 2.1.4). We use *word2vec* to perform the conversion. As each entity and relation is a word, we test the similarity of every word in the question to every entity and relation. The two highest scoring graph labels (at least one entity) are the knowns in our query. If there are two entities, then we have a link prediction query. If there is an entity and a relation, then we have an entity prediction query. The advantage of using word2vec similarity is that SAT questions often rephrase the textual words. For example, if in the passage a character is “angry”, the question will use a word such as “incensed”, which does not exist in the knowledge graph, instead. word2vec translates unseen words into knowledge graph labels which can then be operated on.

Link-prediction queries are relatively simple to answer. Given e_i, e_j we return the relation k that satisfies $\arg \max_k \hat{\mathcal{X}}_{ijk}$. This is the relation which our factorization guesses is strongest between the two entities.

Entity prediction queries are answered similarly. Given e_i, k we return the entity e_j that satisfies $\arg \max_j \hat{\mathcal{X}}_{ijk}$. This is the entity which our factorization guesses is most likely to be the recipient of relation k from entity e_i .

Once we find the label that satisfies the query, we test the similarity of this label to every word in every answer choice, and pick the answer choice that matches best. To summarize, we first convert the question to a query on the knowledge graph, find the label in the knowledge graph that answers the query, and select the answer choice that is most similar to it.

4.3 Memory Networks

The second way we propose answering queries is with neural networks - specifically, an enhanced version of Memory Networks (also referred to as *MemN2N*). The advantage of the neural network approach is that it is very intuitive, computationally feasible, and empirically proven. However, it does lack the strong theory of the mathematical approach. We choose the Memory Networks architecture as the basis for our model because it has strong performance on the bAbi dataset and the foundational architecture itself scales easily to more complex problems as long as the input is in a valid format. First we review the architecture as proposed in [25]. Then we discuss how to adapt it to this specific problem.

4.3.1 Memory Network Architecture

First we will review the architecture for one layer, as the extension to several layers is straightforward. The basic model takes in a set of input vectors x_1, \dots, x_n and a query vector q , and outputs an answer a . Memory Networks require the input text and the query to be vectorized in some form for easy computation. These embedding techniques are covered in detail in Section 2.1. Typically the embedding is done with bag-of-words embedding. This embedding technique relies on the same word being used in the answer choices as well as in the original text, which is true for MCTest and bAbi but not for SAT questions. The vocabulary comes from a dictionary of fixed size V .

The input vectors are converted into a set of memory vectors m_i using a matrix A of size $d \times V$, where d is a pre-determined dimension. q is also embedded into a state u using a

different matrix B of size $d \times V$. To find which memory vectors are most relevant to the initial query, we take the dot product of each vector with u .

$$p_i = \text{Softmax}(u^T m_i)$$

Since the softmax function returns a value between 0 and 1, each score p_i is a scalar that can be considered a probability of how relevant each memory vector is to the query. The more relevant it is, the more likely it is to contain the answer.

The input vectors are now sent through another matrix C to produce a set of transformed inputs c_i . Each transformed vector c_i is weighted by its probability p_i and the sum is taken as the output o of that layer.

$$o = \sum_i p_i c_i$$

This output can then be fed as an input to the next layer, but in the single layer case it is added to u and sent through one final matrix W of size $V \times d$. Lastly, a softmax is taken on the result to yield a probability distribution over the vocabulary space. This is designed to pick one word or a set of words as the answer.

$$\hat{a} = \text{Softmax}(W(o + u))$$

The network itself is trained using backpropagation. The answers to the training set questions are provided as input, and the network learns to answer them.

Several strategies can be used to extend the Memory Network to multiple layers. We use the *Recurrent Neural Network* strategy. This strategy is the simplest to understand as well as to train. Each layer acts the same as in the one-layer case. The only difference is that for all layers $k > 1$ the input $u^{k+1} = u^k + o^k$, and the matrix W is only applied to o^{final} . Each layer uses the same A, B, C matrices. By using the same matrices at each level and passing the state along, the architecture acts like a Recurrent Neural Network.

The initial layer calculates how relevant each memory vector is to the initial query. That output is combined with the initial query and passed as u to the next layer. That layer finds how relevant each memory vector is to the new state u . This means that memories which were not relevant to the initial query may be relevant to the memories which were relevant to the initial query. Each layer finds new memories that can help answer the question.

4.3.2 Memory Network Enhancements

As discussed, the original Memory Networks model uses simple bag-of-words embedding to vectorize the text into the input vectors x_i . The problem with this method is that the “most relevant” memories found in each layer are relevant based on the exact word choice, not based on conceptual similarity. To answer SAT questions the text must be vectorized in such a way that the knowledge graph relationships are captured, and the question must be converted into a query vector for the network. We need to change the inputs while preserving the overall architecture so that the memories chosen as relevant actually match the concept in the query, as opposed to simply matching the words.

We propose stacking each slice of the tensor \mathcal{X} (Section 4.2.1) horizontally and using each column as an input vector. In total the input matrix \mathcal{X}' will be of shape $n \times (nm)$. The query is more complicated to vectorize. The memory network architecture requires the query to be a 1-dimensional vector. We create a vector q of length $N + M$ where N is the number of entities and M the number of relations and set each element to 0. We find the two most similar labels as in Section 4.2.5 and set their corresponding vector elements to 1. This is a two-hot encoding vectorization.

We make two significant changes to the architecture. We set the current layer $u^k = R(u^{k-1} + o^{k-1})$, and define o as the dot product of Cx_i and p vector instead of the scaled sum. Adding another matrix R allows us to capture extra complexity across layers while preserving dimensionality. Changing o to be a dot product is also to help with dimensionality. The original matrix had a fixed query size as well as input length of 177, but with SAT

questions the length $N + R$ is different for every knowledge graph input, so we need to add more flexibility in the architecture for this changing dimensionality. We also update the dimensions of the other matrices so that they handle the length N input vectors, length $N + M$ query vector, and length 5 output vectors.

The architecture equations become

$$\begin{aligned}
m_i &= Ax_i \\
u^0 &= Bq, u^{k+1} = R(u^k + o^k) \\
p_i &= \text{Softmax} \left(\left(u^k \right)^T m_i \right) \\
c_i &= Cx_i \\
o^k &= \sum_i p_i c_i \\
\hat{a} &= \text{Softmax} \left(W \left(o^l + u^l \right) \right)
\end{aligned}$$

where l is the total number of layers.

In the multiple-layer case, the approach is the same as in the original model where the matrix C stays constant at each layer, and $u^{k+1} = u^k + o^k$. We use a 10-layer neural network. The original paper used a 3-layer model to find 3 sets of relevant facts. SAT questions require information pieced together from many sentences, however, so we add more layers to capture this additional complexity. Our architecture is trained with backpropagation.

4.4 Model Comparison

The Tensor model and the Memory Network model have the same underlying goal but achieve it very differently. Both models try to find the inherent properties of the knowledge graph. Both models can use this latent representation of the graph to answer queries about it.

The tensor model has a much stronger mathematical base. The RESCAL model uses ASALSAN [1] which is a general mathematical model for tensor evaluation. Using ALS to optimize is an approximation to make the problem computationally feasible, but the technique is still mathematically sound. At the same time, however, the tensor model is trained only on the knowledge graph itself and not on the answers. It is heavily reliant on the thoroughness of the knowledge graph triples extracted from the text. Information that is missed in the triple extraction will be lost to the tensor model as well.

The Memory Networks model is the exact opposite. It has a much weaker theoretical base but is also more robust to errors in the knowledge graph extraction process. The mathematics of neural networks are not as well understood as that of tensors, especially as we add more layers. The neural net model can behave in various unexpected ways. Nevertheless, the model is trained using the answers to the training set questions. It can find patterns and information using the training process, and as a result is less subject to fluctuations in the actual knowledge graph extraction.

5 Data

The model is evaluated on SAT tests in the old (pre-2016) format. There are several reasons we use this format. The main reason is that the old reading section consists of mostly reading comprehension, while in the new format there are charts and graphs that must be interpreted as well. Reading these requires some computer vision work on top of our existing models, and the questions do not naturally make sense as queries on a knowledge graph. Also, the old format was tried and tested for many years while the new format is just in its first year. Additionally, there are far more resources for the old SAT which we can use as data.

There are no publicly available datasets for the SAT. We tried contacting ETS to obtain a research dataset, but as we have not yet heard a response, we collected practice SAT tests from a 2005 CollegeBoard preparation book which was already purchased several years

ago. We evaluated on all reading comprehension sections from 3 full SAT tests, where each practice test contains 3 reading sections. For convenience, we use “test” to refer to just the 3 reading sections of a test, and “reading section” to refer to just the reading comprehension passages and accompanying questions of each reading section.

Within each section there can be *small* passages (about 100 words) with 4 questions, *medium* passages of 450-500 words and 5-8 accompanying questions, and *long* passages of 500-600 words with 12-13 questions. Of the first two sections, each has one small passage and either 2 medium passages or 1 long passage. The third reading section always has only one long passage. Each question has an associated difficulty as well (*easy*, *medium*, or *hard*). The Math and Writing sections, as well as the vocabulary questions in the reading section, are not used in evaluation.

Our practice tests were initially in paper format and needed to be converted to an electronic format. This conversion was actually quite complicated. First, every test was scanned into PDF format. Then these PDFs were converted into text format using a combination of Amazon Mechanical Turk and OCR software. For Mechanical Turk we split each test into individual passages - 2 small, 2 medium, and 2 long. Each passage was published as its own task for workers. The PDFs were converted into images and the images for each passage plus questions were uploaded to Amazon Web Services S3 File Storage. Each Mechanical Turk task placed the images on the left side of the screen and text boxes to type the passage on the right side of the screen. With the OCR software, copies of each image had to be made. Each SAT page has two columns, which are misinterpreted by the OCR software as a single line, so the right-hand side and left-hand side of the image were processed separately. Tests 2 and 3 were processed via Mechanical Turk while Test 1 was processed via OCR software. The expenditure on Mechanical Turk was \$116.92 and expenditure on OCR was \$17.95, for a total of \$134.87. Due to some errors in the conversion process, a few questions (mainly on small passages) are not used in the full evaluation. The passage text, question text, and answer key are all stored in text form and used as input for evaluation (Section 6.2).

6 Evaluation and Discussion

In this section we present and interpret our results on three full SAT reading comprehension tests. First we describe the implementation details, evaluation process, and raw results. Then we compare the performance of our enhanced models to their baseline versions and discuss the performance of each model on each test. We break down what types of problems each model did well / poorly on, and offer an interpretation as to why. Finally we compare the Tensor model to the Memory Networks model and evaluate their relative performance.

6.1 Implementation Details

There are three main components to the implementation: knowledge graph creation, tensor decomposition, and neural network. We heavily modified the Java code provided by Narasimhan and Barzilay¹ [19], which itself uses the Stanford CoreNLP Toolkit² [17], to create the entity-relation triple extraction program. We modify the RESCAL algorithm created by Nickel³ [20] to include our customized regularization term. The memory network is implemented using Numpy⁴, Theano⁵, and Jupyter Notebook⁶ [14, 22, 2, 4].

6.2 Evaluation Process

The evaluation process starts with the text input separated into passage text, questions, and answers. For each passage, the passage text is fed through the knowledge extraction program, which converts it into a set of entity-relation triples and outputs those into another text file.

The tests are answered passage-by-passage. This means that each model is evaluated one passage at a time. Each passage is considered its own knowledge graph, so in the

¹<https://bitbucket.org/ghostof2007/mctest>

²<http://stanfordnlp.github.io/CoreNLP/>

³<https://github.com/mnick/rescal.py>

⁴<http://www.numpy.org/>

⁵<http://www.deeplearning.net/software/theano/>

⁶<http://ipython.org/notebook.html>

tensor model a new tensor of varying dimensionality is created for each passage, and the accompanying questions can only be answered as queries on that particular graph. Combining the knowledge graphs for different sections would only serve to make it harder to answer the questions (and it is already hard enough!) Therefore it makes sense to run the tensor model on each passage separately. The same holds true for the Memory Networks model. The neural net has an additional training constraint. Due to the dimensionality constraints on the embedding matrices and the underlying architecture, it is not possible to combine the tensors from different passages.

In each evaluation we first train the model and then evaluate its performance on the questions for that passage. The Tensor model is trained using ALS (Section 4.2.4). The training process is independent of the actual performance of the model, i.e. there is no feedback, so we can test the model’s performance on every question. The neural net is trained in a feedback loop using backpropagation as discussed in Section 2.2.1 and Section 2.2.2. In each passage we set aside $\frac{2}{3}$ of the questions for training and use the remaining $\frac{1}{3}$ as the test set. For this reason we do not evaluate the neural net on any of the Small passages.

We run the Tensor model 10 times with $d = 20$ and pick the model with best performance. We train the Neural Network model for 1000 iterations, evaluating its performance every 50 iterations, and pick the model with best performance. Ideally we would have a validation set separate from the test set, but because we have to train passage-by-passage the number of questions is already quite low for each model.

6.3 Results

Appendix Section 9.4 contains full evaluation data for each test. The tables contain each model’s answer for every question, as well as the correct answer for that question and its associated difficulty. Cells are marked in red if the model answered the question correctly. We evaluated Test 1 on our enhanced models as well as their baseline versions for comparison. Both models outperformed the baselines, so on Tests 2 and 3 we evaluated only on our

enhanced models. The analysis of exactly how our models perform better can be found in Section 6.4.

On Test 1 our Tensor model with regularization scores 8 / 41 while the model without regularization scores 7 / 41. The 1-layer and 10-layer Memory Networks both scored 2 / 13. Hereafter “tensor model” refers to the tensor model with regularization, and “MemN2N” refers to the 10-layer neural network.

On Test 2 the tensor model scored 11 / 40 while the MemN2N scored 7 / 14.

On Test 3 the tensor model scored 9 / 41 while the MemN2N scored 3 / 14.

6.4 Comparison to Baseline

On Test 1 we compared our enhanced model to their respective baselines. We updated the tensor model with SSE regularization, so its baseline is the model without regularization. The baseline for the 10-layer MemN2N is the 1-layer MemN2N.

The tensor model with regularization scores 1 better than the model without regularization. Besides pure score, there is additional evidence that adding regularization creates a better model. For each passage we train the model 10 times and evaluate its performance. Without regularization, the training process produces the same model each time. However, with regularization it produces a slightly different model each time. The optimization process with ALS approximates the graph structure, which means that the tensor factorization should be slightly different each time. If it is the same, as is the case in the model without regularization, then the inherent properties are being found *too easily*, suggesting that it is not a good factorization of the graph.

The regularization term we use imposes some geometric constraints on the model. Training without regularization allows the training to violate these constraints, producing a model which does not actually represent the underlying graph. Adding the regularization term keeps the model robust by forcing the factorization to more accurately represent the underlying graph. The model is not only more likely to answer correctly, but also more resistant to

errors in the knowledge graph extraction. We discussed in Section 4.4 that the tensor model is too reliant on the knowledge graph extraction process. Ensuring the model is robust is of utmost importance, and the regularization term is necessary to do that.

With regards to the neural net, the 10-layer MemN2N is significantly better than the 1-layer version. Although both score the same on Test 1, the 1-layer MemN2N does not actually learn any of the underlying representation! Table 1 demonstrates that in each section, the 1-layer MemN2N predicted the same answer choice for every question. Its choice was the letter answer choice seen most in the training set. For example, in the Section 2 first Medium passage there are two questions where the answer is *C*. No other answer choice appears more than once. As a result, the 1-layer MemN2N guesses *C* for all of the remaining questions. The 10-layer MemN2N, on the other hand, guesses different answers for different questions in the same section. This indicates that it is predicting an answer based on some underlying semantic information, rather than simply the most commonly-seen answer choice. Simply guessing the most commonly-seen answer choice is not useful in evaluating how well the model understands the semantic information of the questions and the underlying knowledge graph. The reason the 1-layer MemN2N only outputs the most commonly-seen answer choice is that it overfits very quickly. We can guard against overfitting and create a more flexible architecture by adding more layers. Therefore the 10-layer MemN2N is our model of choice.

6.5 Interpretation of Results

In this section we interpret the models' performance. We first note that the model performance is not consistent across passages, but rather very good on some passages and very poor on others. For example, on Test 2 Section 1 the tensor model scores 3 / 12 but on Test 2 Section 2 on the second Medium passage the tensor model scores 4 / 9. Similarly, on Test 2 Section 3 the MemN2N scores 4 / 5 but on Test 3 Section 2 the MemN2N scores 0 / 4.

There are several possible reasons for this phenomenon. The first is that both models are dependent on the knowledge graph extraction, so if the extraction process is accurate for a given passage then both models will perform well, and if the extraction is inaccurate then both models will perform poorly. This reason is unlikely, however, because the two models do not necessarily perform well or poorly in the same passages. On Test 2 Section 3 the MemN2N scores 4 / 5 but the tensor scores just 2 / 13. If the knowledge graph extraction were influencing the consistency of our models, they should both perform equally well or poorly in a given section.

The more likely reason, and the one that proves the usefulness of our models, is that within those passages the models are able to find the inherent properties of the knowledge graph (and therefore the text) and answer questions about it. Why the models work on some passages but not on others is unclear, but that is to be expected when working with machine learning algorithms. While the mathematics behind the tensor method are clear, the actual ALS process is still quite variable and hard to interpret. The neural network is even harder to understand, as its learning process is hidden behind 10 layers. However, we know that when the models are scoring very well on multiple sections it is not simply coincidence. That they score very poorly on multiple sections as well simply reinforces the idea that the models are answering questions based on their understanding of the text. Whether or not they can answer correctly depends on how well they parse the question and how accurate their representation of the text is. What dictates this accuracy is not easily seen due to the learned nature of the models. The goal of this research was to create models which would be able to actually answer questions based on the semantic information, instead of just using bag-of-words matching or other basic techniques, and we have done just that.

This claim is strengthened when we investigate the passages which the models performed particularly poorly on. Test 1 Section 3 has perhaps the hardest passage for our models, with the tensor model scoring 0 / 13 and the MemN2N scoring 1 / 5. This passage is actually composed of two separate passages about the same topic, and most of the questions require

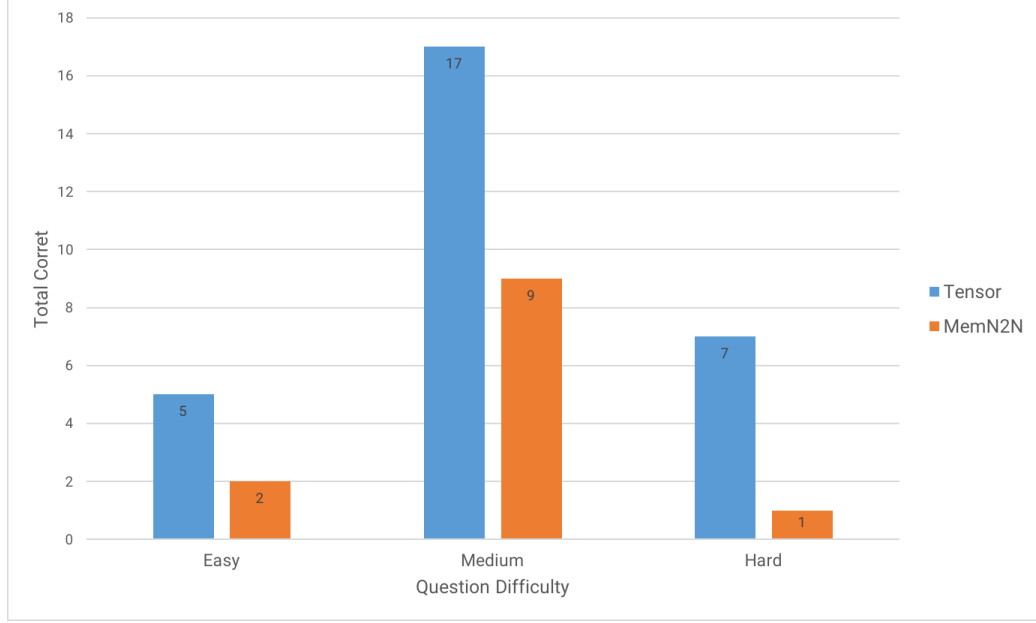


Figure 5: Correct answers by difficulty level

comparing the opinions and viewpoints of the two authors. These types of comparison questions cannot be answered using the knowledge graph, because they require a higher level of analysis. Our models try to answer the question by completing a triple using the entities and relations on the knowledge graph, but the types of questions asked in Test 1 Section 3 require analyzing the entire knowledge graph itself. We expect our models to score poorly on such questions, and that is what happens. It is difficult to explain exactly what about certain passages makes our models score well, but we can understand which passages will cause our models to score poorly.

6.6 Performance by Difficulty

Figure 5 shows the total number of questions answered across all three tests sorted by their difficulty, while Figure 6 uses the same data but controlled for the total number of questions of each difficulty.

The first observation we can make is that both models answer significantly more Medium questions than Easy or Hard, but that their percentage on Medium questions is lowest.




Figure 6: Percent of questions of each difficulty level answered correctly

This point is actually quite subtle. The full data tables demonstrate that there are far more Medium questions in general than Easy or Hard. However, the huge number of Medium questions means that the types of Medium questions are also quite varied. Whether the percentage of Medium questions answered correctly is so low is because there are simply far more Medium questions or because there are types of Medium questions that both models score poorly on (such as compare/contrast questions across passages) is unclear.

The second observation we make is that the MemN2N model scores perfectly on the Easy questions and 50% on the Hard questions. The problem with analyzing this point is that the sample size is simply too small to make any claims about the MemN2N’s performance on Easy or Hard questions. We can definitively say that the model was able to answer at least one of each type of question, even given the limited test size, but we cannot claim more than that.

The third important observation here is that the tensor model answers a very high percentage of Hard questions correctly. This is partially due to the relatively small sample size (only 11), but also a consequence of the way in which the tensor model is actually answering

these questions. Consider Test 3 Section 2, where the tensor model answers 2 out of 4 Hard questions correctly. This is actually another passage that contains two smaller passages, but both of the Hard questions the model gets right deal with only one of the passages. In fact, both questions ask about the meaning of a specific phrase in the context of the passage. The tensor model is actually able to match the words in the phrase with entities in the knowledge graph, and then select the answer choice that best corresponds to the matched entity.

As far as questions go, these are actually some of the easiest for the tensor model to answer because they do not require any high-level analysis. They simply require matching words in the question to entities in the graph, and then matching the corresponding graph element to words in the answer choice. This is exactly how the tensor model works, by using word2vec to match words in the question to entities in the graph, and then matching the relevant information to words in the answer choices.

This analysis leads us to a key observation. The difficulty standards for humans are not at all similar to the difficulty standards of our models. The Hard questions described in the paragraph above can be very difficult for human readers, because they require understanding the given phrase in context and then translating it into one of the answer choices. Essentially, these questions require fully understanding the phrase as well as its context. Our models, however, do not need to fully understand the phrase or its context; they simply need to match words in the question to entities in the graph. On the other hand, questions marked Easy can be very difficult for our models (namely the tensor model, as the MemN2N scored 2 / 2 on all its Easy questions). Some Easy questions ask about the author's possible opinion on some topic that is not directly contained in the text. For many humans these can be answered simply by understanding the author's tone throughout the passage. Our models cannot understand tone, and cannot process information not contained in the knowledge graph. These types of questions are therefore impossible for them to answer.




Figure 7: Percent of questions on each test answered correctly

6.7 Comparing Tensor Model to Memory Networks

Figure 7 shows the percent of questions answered correctly by each model on each test. There are several striking observations. The first is that the tensor model scores quite consistently, varying by just over 10% across all three tests. The MemN2N model is very volatile, in fact scoring the single lowest and highest performances across both models and all tests.

This result is to be expected, as we discussed in Section 4.4 how the tensor model is more mathematically sound while the neural net model does not have as clear mathematics describing how it works. The MemN2N’s flexibility allowed it to score very high on Test 2, but at the same time caused it to score very low on Test 1.

We also note that the models scored relatively similar on all of the tests. They both scored lowest on Test 1, second-best on Test 3, and highest on Test 2. Additionally, on both Test 1 and 3, the scores for the tensor and MemN2N are almost identical, differing by no more than 2%. This result is partially due to the small sample size of MemN2N questions. Nevertheless, that the two models have similar scores across all three tests indicates that

for a given knowledge graph they find the hidden representation of the knowledge graph equally well. We conclude that the two models are, for the most part, equally viable. Across the three tests, the tensor model scores barely higher than the MemN2N model on Tests 1 and 3, but significantly worse on Test 2. The results demonstrate that the tensor model is the better general choice, as it guarantees some amount of competence, but at the same time it will not score very high. The MemN2N will most likely score slightly worse than the tensor model, but has much higher potential.

7 Future Work

We hope this work is just the start of research into answering SAT questions and tackling machine comprehension tasks of similar difficulty. To that end, we present several ways to improve the models discussed in this paper. These improvements fall in three major areas: linguistic analysis, tensor model, and neural network.

7.1 Linguistic Analysis

There are three instances in our model where we need to analyze the semantic or structural information of text. The first is extracting the knowledge graph triples from the passage text. The second is parsing the question and understanding what it asks for. The third is choosing an answer choice given the model output - this applies only to the tensor model.

7.1.1 Knowledge Graph Extraction Improvements

There are some drawbacks to the way we currently extract the knowledge graph triples from the text. Some information is lost in the extraction process. Because the subject is forced to be *entity₁* we lose information about the direct object. For example, in the sentence in Figure 3, our models learn that the dog likes sausage but not that the sausage is spicy.

Additionally, SAT sentences are very difficult to parse. Most have multiple clauses, and normally unusual structural elements such as the colon or dash are quite commonplace. One sentence used in an actual passage is: *The critic Edmund Wilson was not a self-conscious letter writer or one who tried to sustain studied mannerisms.* We are able to extract the triple $\langle \text{Wilson, was, writer} \rangle$ but due to the complicated structure of the sentence we are not able to extract that Wilson was not self-conscious or that Wilson did not “try to sustain studied mannerisms.” Overall we can extract the most important points, but we often miss the subtleties of the sentence.

If we can add more semantic rules to triple extraction then both the tensor model as well as the MemN2N could have more information at hand to answer the question. Right now we only focus on the *nsubj* and *dobj* or *cop* dependencies. One useful dependency we could add is the *amod* dependency, which is for adjectival modifiers [10]. This could help us capture descriptive information about entities that is not given with a copular verb. We can also focus on entities that are not the nominal subject, or add logic to handle sentences with multiple clauses.

7.1.2 Question Parsing

In our existing model, we vectorize a given question by using word2vec to calculate the similarity of every word in the sentence (excluding stopwords) to every possible entity and relation, and choosing the 2 entities or entity and relation that have highest similarity to question words. This works well when the question tests a specific concept, but not when the question is more complicated. This system runs into the same problem as bag-of-words matching, where the semantic and structural information of the sentence is lost.

A better way to tackle question parsing would be to feed the questions through the knowledge graph extraction program, and extract the information from the question. However, because we need to frame the question as a query on the knowledge graph we would modify the extraction program to extract triples with a missing entity or relation. We can then try to

fill in these missing elements using the original knowledge graph, and match the resulting filled-in triples to the answer choices. The answer would be the choice with strongest match. Implementing this would require the improved deep extractor discussed above in Section 7.1.1.

7.1.3 Choosing an Answer

Currently in the tensor model we choose an answer similarly to the way we parse a question- by using word2vec to find the single word which matches the model’s output best. Often, the answer choices are sentences too, so we need to be able to parse them. Just as with the questions we can feed the choices through the Knowledge Graph extractor to find a set of triples. We then need to create some metric to match triples and use it to pick the best choice.

7.2 Tensor Model Improvements

There are several improvements that can be made to the tensor model. The first is that the model’s accuracy can be affected greatly by errors in the knowledge graph extraction process. We need to find a way to make the tensor model robust to such fluctuations. The best way to do this is to use an ensemble training method. We can take the knowledge graph tensor and create several variants, each slightly different from the original. Each of these are trained separately, and then combined to create the actual approximation. This approximation has several redundancies and is therefore less affected by errors in the original knowledge graph. Ensemble learning methods are common in many machine learning algorithms.

Another improvement that can be made to the tensor model is to provide supervision in the training process. Because the model is not trained on the actual questions, it cannot learn from the structure of the questions. By adding a small supervised training component, we can achieve the flexibility of the MemN2N while keeping the performance guarantee of the tensor model.

7.3 Memory Network Improvements

There are two main improvements we can make to the Memory Network. The first is with the underlying architecture itself. The input vectors must be embedded with matrices A, C . Since each passage has a tensor of different shape, we cannot reuse A, C across passages. We discussed in Section 6.2 how it is advantageous to train on just one passage at a time. However, with the neural network there may be some underlying patterns across passages that we are not able to capture currently. If the dimensionality constraints on the architecture are loosened, the MemN2N model can be trained across passages and even across tests. This means that the model could be trained on one particular test, and run fully on another test.

The second possibility is to improve the model's consistency while preserving its flexibility. The biggest drawback of the MemN2N currently is that it occasionally has very poor performance. If we can train the model such that it will be able to answer certain types of questions with high probability then the model will be more reliable. For example, the tensor model is able to answer most questions that ask for the meaning of some phrase in context (Section 6.6). If the MemN2N model is more reliable and consistent while still flexible, then it can become the best tool for machine comprehension. The very nature of neural network training makes it hard to create consistency, but if the MemN2N has the ability to train across tests then a large training corpus could help it achieve the desired consistency.

8 Conclusion

In this work we present an approach to machine comprehension tasks that uses semantic and structural information instead of simple word-matching. We discuss how any text can be represented as a knowledge graph of entity-relation triples and questions about the text can be converted into queries on the graph. These queries can be answered with either link prediction or entity prediction. We use the Stanford NLP Toolkit's Dependency Parser

[17, 9] to turn every sentence in the text into a dependency graph, and then use semantic rules to extract the information in triple form and create the knowledge graph. We present two methods to answer the knowledge graph queries. We first present a tensor decomposition approach, which factorizes the knowledge graph into a product of matrices representing the latent properties of the graph. We improve the RESCAL [20, 8] method to use Semantically Smooth Embedding [11]. The second approach uses neural networks. We enhance the Memory Networks [28, 25] architecture and generalize it so that it can take any arbitrary knowledge graph and question as input.

We evaluate the tensor decomposition and Memory Networks methods on three full SAT reading comprehension tests. We demonstrate that the methods are able to answer questions correctly using the actual semantic and structural meaning of the original text. We also prove that our enhanced models perform better than their respective baselines. The tensor decomposition model mostly outperforms the Memory Networks model, but the Memory Networks model has a higher potential to perform well. We discuss the types of questions which are possible for these models to answer and which require a higher level of analysis. Finally, we present possible extensions to our models. These include stronger semantic parsing, making the tensor model more robust, and improving the Memory Network’s consistency.

The main goal of this research was to create a general model that could be used to tackle a machine comprehension task that is much harder than the standard MCTest [23] and bAbi [27] datasets. We wanted to show that these models could actually capture the semantic and structural information of the text, and use that to answer questions. In this regard we have definitely succeeded. Both models are able to parse the meaning of the text as well as the questions, and use that information to find the correct answers. There is still much room for improvement, however. We hope that the ideas presented here will be useful in further machine comprehension research.

References

- [1] B. Bader, R. a. Harshman, and T. G. Kolda, “Temporal Analysis of Semantic Graphs Using ASALSAN,” *Seventh IEEE International Conference on Data Mining (ICDM 2007)*, pp. 33–42, 2007. [Online]. Available: <http://ieeexplore.ieee.org/lpdocs/epic03/wrapper.htm?arnumber=4470227>
- [2] F. Bastien, P. Lamblin, R. Pascanu, J. Bergstra, I. J. Goodfellow, A. Bergeron, N. Bouchard, and Y. Bengio, “Theano: new features and speed improvements,” *Deep Learning and Unsupervised Feature Learning NIPS 2012 Workshop*, 2012.
- [3] J. Berant and P. Clark, “Modeling Biological Processes for Reading Comprehension,” *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 1499–1510, 2014. [Online]. Available: <http://allenai.org/content/publications/berant-srikumar-manning-emnlp14.pdf>
- [4] J. Bergstra, O. Breuleux, F. Bastien, P. Lamblin, R. Pascanu, G. Desjardins, J. Turian, D. Warde-Farley, and Y. Bengio, “Theano: a CPU and GPU math expression compiler,” in *Proceedings of the Python for Scientific Computing Conference (SciPy)*, Jun. 2010, oral Presentation.
- [5] C. M. Bishop, *Neural networks for pattern recognition*, 1995, vol. 92.
- [6] A. Bordes, N. Usunier, A. Garcia-Duran, J. Weston, and O. Yakhnenko, “Translating embeddings for modeling multi-relational data,” in *Advances in Neural Information Processing Systems*, 2013, pp. 2787–2795.
- [7] D. Britz. (2015) Introduction to rnns. [Online]. Available: <http://www.wildml.com/2015/09/recurrent-neural-networks-tutorial-part-1-introduction-to-rnns/>
- [8] K.-W. Chang, W.-t. Yih, B. Yang, and C. Meek, “Typed tensor decomposition of knowledge bases for relation extraction,” *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 1568–1579, 2014.
- [9] D. Chen and C. D. Manning, “A Fast and Accurate Dependency Parser using Neural Networks,” *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, no. i, pp. 740–750, 2014. [Online]. Available: <https://cs.stanford.edu/~danqi/papers/emnlp2014.pdf>
- [10] M.-C. De Marneffe, T. Dozat, N. Silveira, K. Haverinen, F. Ginter, J. Nivre, and C. D. Manning, “Universal stanford dependencies: A cross-linguistic typology,” in *LREC*, vol. 14, 2014, pp. 4585–4592.
- [11] S. Guo, Q. Wang, B. Wang, L. Wang, and L. Guo, “Semantically smooth knowledge graph embedding,” in *Proceedings of ACL*, 2015, pp. 84–94.
- [12] L. Hirschman, M. Light, E. Breck, and J. D. Burger, “{D}eep {R}ead: {A} Reading Comprehension System,” *Proceedings of ACL*, pp. 325–332, 1999.
- [13] S. Hochreiter, “Untersuchungen zu dynamischen neuronalen netzen,” *Master’s thesis, Institut für Informatik, Technische Universität, München*, 1991.
- [14] E. Jones, T. Oliphant, P. Peterson *et al.*, “SciPy: Open source scientific tools for Python,” 2001–, [Online; accessed 2016-04-23]. [Online]. Available: <http://www.scipy.org/>
- [15] Y. Koren, R. Bell, and C. Volinsky, “Matrix Factorization Techniques for Recommender Systems,” *Computer*, vol. 42, no. 8, pp. 42–49, 2009.
- [16] A. Kumar, O. Irsoy, J. Su, J. Bradbury, R. English, B. Pierce, P. Ondruska, I. Gulrajani, and R. Socher, “Ask Me Anything: Dynamic Memory Networks for Natural Language Processing,” *arXiv*, pp. 1–10, 2015.
- [17] C. D. Manning, M. Surdeanu, J. Bauer, J. Finkel, S. J. Bethard, and D. McClosky, “The Stanford CoreNLP Natural Language Processing Toolkit,” *Proceedings of 52nd Annual Meeting of the ACL: System Demonstrations*, pp. 55–60, 2014.
- [18] T. Mikolov, K. Chen, G. Corrado, and J. Dean, “Distributed Representations of Words and Phrases and their Compositionality,” *Nips*, pp. 1–9, 2013.
- [19] K. Narasimhan and R. Barzilay, “Machine Comprehension with Discourse Relations,” *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pp. 1253–1262, 2015. [Online]. Available: <http://www.aclweb.org/anthology/P15-1121>
- [20] M. Nickel, V. Tresp, and H.-P. Kriegel, “A Three-Way Model for Collective Learning on Multi-Relational Data,” *28th International Conference on Machine Learning*, pp. 809–816, 2011.
- [21] M. A. Nielsen, *Neural Networks and Deep Learning*. Determination Press, 2015. [Online]. Available: <http://neuralnetworksanddeeplearning.com/index.html>
- [22] F. Perez and B. E. Granger, “Ipython: A system for interactive scientific computing,” *Computing in Science Engineering*, vol. 9, no. 3, pp. 21–29, May 2007.
- [23] M. Richardson, C. J. C. Burges, and E. Renshaw, “MCTest: A Challenge Dataset for the Open-Domain Machine Comprehension of Text,” *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing (EMNLP 2013)*, no. October, pp. 193–203, 2013.

- [24] S. T. Roweis and L. K. Saul, “Nonlinear dimensionality reduction by locally linear embedding,” *Science*, vol. 290, no. 5500, pp. 2323–2326, 2000.
- [25] S. Sukhbaatar, A. Szlam, J. Weston, and R. Fergus, “End-To-End Memory Networks,” pp. 1–11, 2015. [Online]. Available: <http://arxiv.org/abs/1503.08895>
- [26] Z. Wang, J. Zhang, J. Feng, and Z. Chen, “Knowledge graph embedding by translating on hyperplanes.” in *AAAI*. Citeseer, 2014, pp. 1112–1119.
- [27] J. Weston, A. Bordes, S. Chopra, T. Mikolov, and A. M. Rush, “Towards AI-Complete Question Answering: A Set of Prerequisite Toy Tasks,” 2015. [Online]. Available: <http://arxiv.org/abs/1502.05698>
- [28] J. Weston, S. Chopra, and A. Bordes, “Memory Networks,” *International Conference on Learning Representations*, pp. 1–14, 2015. [Online]. Available: <http://arxiv.org/abs/1410.3916>

9 Appendix

9.1 MCTest Example

Figure 8 shows a sample MCTest passage and an accompanying question. The sentence structure is straightforward throughout the passage, and the question can be answered using bag-of-words matching.

James the Turtle was always getting in trouble. Sometimes he'd reach into the freezer and empty out all the food. Other times he'd sled on the deck and get a splinter. His aunt Jane tried as hard as she could to keep him out of trouble, but he was sneaky and got into lots of trouble behind her back.

One day, James thought he would go into town and see what kind of trouble he could get into. He went to the grocery store and pulled all the pudding off the shelves and ate two jars. Then he walked to the fast food restaurant and ordered 15 bags of fries. He didn't pay, and instead headed home.

His aunt was waiting for him in his room. She told James that she loved him, but he would have to start acting like a well-behaved turtle.

After about a month, and after getting into lots of trouble, James finally made up his mind to be a better turtle.

1) What is the name of the trouble making turtle?

- A) Fries
- B) Pudding
- C) James
- D) Jane

Figure 8: A sample MCTest passage and accompanying question, from [23]

9.2 SAT Example

Figure 9 shows a sample SAT reading comprehension passage and some accompanying questions. The sentence structure here is far more complex than in the MCTest passages. The vocabulary used is more varied and nuanced as well. Additionally, the questions are tougher to parse and require information across multiple sentences and even paragraphs. The answer choices are harder to interpret as well.

Questions 10-15 are based on the following passage.

This excerpt from a novel by a Chinese American author is about a Chinese American woman named June. During a family dinner party attended by some of June's Chinese American friends, Waverly, a tax attorney, discusses an advertisement that June wrote for her.

Waverly laughed in a lighthearted way. "I mean, really, June." And then she started in a deep television-announcer voice: "Three benefits, three needs, three reasons to buy . . . Satisfaction guaranteed . . ."

Line 5 She said this in such a funny way that everybody thought it was a good joke and laughed. And then, to make matters worse, I heard my mother saying to Waverly: "True, one can't teach style. June is not sophisticated like you. She must have been born this way."

10 I was surprised at myself, how humiliated I felt. I had been outsmarted by Waverly once again, and now betrayed by my own mother.

.....

Five months ago, some time after the dinner, my mother gave me my "life's importance," a jade pendant on a gold chain. The pendant was not a piece of jewelry I would have chosen for myself. It was almost the size of my little finger, a mottled green and white color, intricately carved. To me, the whole effect looked wrong: too large, too green, too garishly ornate. I stuffed the necklace in my lacquer box and forgot about it.

15 But these days, I think about my life's importance. I wonder what it means, because my mother died three months ago, six days before my thirty-sixth birthday. And she's the only person I could have asked to tell me about life's importance, to help me understand my grief.

20 I now wear that pendant every day. I think the carvings mean something, because shapes and details, which I never seem to notice until after they're pointed out to me, always mean something to Chinese people. I know I could ask Auntie Lindo, Auntie An-mei, or other Chinese friends, but I also know they would tell me a meaning that is different from what my mother intended. What if they tell me this curving line branching into three oval shapes is a pomegranate and that my mother was wishing me fertility and posterity? What if my mother really meant the carvings were a branch of pears to give me purity and honesty?

25 And because I think about this all the time, I always notice other people wearing these same jade pendants—not the flat rectangular medallions or the round white ones with holes in the middle but ones like mine, a two-inch oblong of bright apple green. It's as though we were all sworn to the same secret covenant, so secret we don't even know what we belong to. Last weekend, for example, I saw a bartender wearing one. As I fingered mine, I asked him, "Where'd you get yours?"

30

35

40

45

"My mother gave it to me," he said.
I asked him why, which is a nosy question that only one Chinese person can ask another; in a crowd of Caucasians, two Chinese people are already like family.
"She gave it to me after I got divorced. I guess my mother's telling me I'm still worth something."
And I knew by the wonder in his voice that he had no idea what the pendant really meant.

50

10. In lines 1-4, Waverly characterizes June's advertisement as being

(A) unsophisticated and heavy-handed
(B) somber and convoluted
(C) clear and concise
(D) humorous and effective
(E) clever and lively

11. In the context of the passage, the statement "I was surprised at myself" (line 10) suggests that June

(A) had been unaware of the extent of her emotional vulnerability
(B) was exasperated that she allowed Waverly to embarrass her in public
(C) was amazed that she could dislike anyone so much
(D) had not realized that her mother admired her friend Waverly
(E) felt guilty about how much she resented her own mother

12. June's observation in lines 10-11 ("I had . . . again") suggests that

(A) June had expected Waverly to insult her
(B) June had hoped to embarrass Waverly this time
(C) Waverly had a private understanding with June's mother
(D) Waverly had made June feel inadequate on previous occasions
(E) Waverly was a more talented writer than June was

13. For June, a significant aspect of what happened at the dinner party is that

(A) her mother had taken great pains to make Waverly feel welcome
(B) her mother had criticized her for arguing with Waverly
(C) her mother had sided against her in front of family and friends
(D) Waverly had angered June's mother
(E) Waverly had lied to June's mother

Figure 9: A sample SAT reading comprehension passage and accompanying questions

9.3 Derivation of Output Layer Error

In this section we present the derivation for the error of an output node of an RNN with cross-entropy loss and softmax at the output layer. Softmax and cross-entropy loss C are defined in Section 2.2.2. Let us consider an output layer with n nodes. Define z_i as the quantity before activation function at the output layer, and \hat{y}_i as the final output of node i . The true output of node i is y_i . The output vectors can be written as $\hat{\mathbf{y}}$, \mathbf{y} . Note that \mathbf{y} is a one-hot vector, so exactly one $y_i = 1$ and all other $y_j = 0$. The output node error is calculated as $\frac{\partial C}{\partial z_i}$.

$$\hat{y}_i = \text{Softmax}(z_i) = \frac{e^{z_i}}{\sum_{j=1}^n e^{z_j}} = \frac{e^{z_i}}{e^{z_i} + \sum_{j \neq i} e^{z_j}}$$

The partial derivatives yield

$$\begin{aligned} \frac{\partial \hat{y}_i}{\partial z_i} &= \frac{\sum_{j=1}^n e^{z_j} (e^{z_i}) - e^{z_i} (e^{z_i})}{\left(\sum_{j=1}^n e^{z_j} \right)^2} = \frac{e^{z_i} \sum_{j=1}^n e^{z_j}}{\left(\sum_{j=1}^n e^{z_j} \right)^2} - \left(\frac{e^{z_i}}{\sum_{j=1}^n e^{z_j}} \right)^2 \\ &= \hat{y}_i - \hat{y}_i^2 = \hat{y}_i (1 - \hat{y}_i) \end{aligned} \tag{3}$$

$$\begin{aligned} \frac{\partial \hat{y}_j}{\partial z_i} &= \frac{\sum_{j=1}^n e^{z_j} (0) - e^{z_j} (e^{z_i})}{\left(\sum_{j=1}^n e^{z_j} \right)^2} = \frac{-e^{z_i}}{\sum_{j=1}^n e^{z_j}} \frac{e^{z_j}}{\sum_{j=1}^n e^{z_j}} \\ &= -\hat{y}_i \hat{y}_j \end{aligned}$$

When we take the partial of C with respect to the individual z_i , we need $\frac{\partial \hat{y}_i}{\partial z_i}$ as well as $\frac{\partial \hat{y}_j}{\partial z_i}$.

$$\frac{\partial \hat{y}_j}{\partial z_i} = \begin{cases} \hat{y}_i(1 - \hat{y}_i) & \text{if } i = j \\ -\hat{y}_i\hat{y}_j & \text{if } i \neq j \end{cases}$$

Now we derive the output node error using the results of Equation (3).

$$\frac{\partial C}{\partial \hat{y}_i} = -\frac{y_i}{\hat{y}_i}$$

$$\frac{\partial C}{\partial z_i} = \sum_{j=1}^n \frac{\partial C}{\partial \hat{y}_j} \frac{\partial \hat{y}_j}{\partial z_i} = \frac{\partial C}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial z_i} + \sum_{j \neq i} \frac{\partial C}{\partial \hat{y}_j} \frac{\partial \hat{y}_j}{\partial z_i}$$

$$\frac{\partial C}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial z_i} = -\frac{y_i}{\hat{y}_i} \hat{y}_i(1 - \hat{y}_i) = -y_i(1 - \hat{y}_i)$$

$$\frac{\partial C}{\partial \hat{y}_j} \frac{\partial \hat{y}_j}{\partial z_i} = -\frac{y_j}{\hat{y}_j} (-\hat{y}_i\hat{y}_j) = y_j\hat{y}_i$$

Summing together these partials yields

$$\frac{\partial C}{\partial z_i} = -y_i(1 - \hat{y}_i) + \sum_{j \neq i} y_j\hat{y}_i = -y_i + y_i\hat{y}_i + \sum_{j \neq i} y_j\hat{y}_i = -y_i + \sum_{j=1}^n y_j\hat{y}_i = -y_i + \hat{y}_i \sum_{j=1}^n y_j$$

However, we know that $\sum_{j=1}^n y_j = 1$ because \mathbf{y} is a one-hot vector. This yields our desired result.

$$\frac{\partial C}{\partial z_i} = \hat{y}_i - y_i$$

This is a very useful result. It can be calculated very easily and in parallel for each node, which prevents backpropagation from becoming a bottleneck and allows the RNN to learn quickly.

9.4 Full Evaluation Data

Here we present the full evaluation data on three practice tests.

Table 1: Evaluation on Test 1

	Number	Difficulty	Correct Answer	Tensor w/o Reg	Tensor w/ Reg	1-layer MemN2N	3-layer MemN2N
<i>Section 1</i>							
Small	9	E	E	A	A	-	-
	10	M	B	A	D	-	-
	11	E	D	D	D	-	-
	12	E	D	E	E	-	-
Long	13	M	E	C	D	train	train
	14	E	D	D	D	train	train
	15	M	B	E	B	train	train
	16	H	A	E	A	train	train
	17	H	A	A	A	train	train
	18	E	C	E	D	train	train
	19	M	A	D	D	train	train
	20	M	B	D	A	train	train
	21	M	B	E	E	A	A
	22	M	C	A	A	A	A
	23	M	B	D	D	A	A
	24	E	A	A	B	A	A
<i>Section 2</i>							
Medium	13	M	C	B	B	train	train
	14	E	B	C	D	train	train
	15	M	C	E	E	train	train
	16	E	E	B	B	train	train
	17	E	D	A	A	C	C
Medium	18	M	C	A	A	C	D
	19	M	E	B	E	train	train
	20	E	C	E	E	train	train
	21	M	D	D	B	train	train
	22	M	E	E	E	train	train
	23	M	A	D	E	E	C
	24	M	B	B	B	E	C
<i>Section 3</i>							
Long	7	M	B	A	A	train	train
	8	M	A	C	C	train	train
	9	M	D	E	E	train	train
	10	M	E	D	D	train	train
	11	M	A	B	B	train	train
	12	E	A	C	C	train	train
	13	E	D	E	B	train	train
	14	M	D	B	B	train	train
	15	M	E	A	A	D	D
	16	M	E	A	A	D	D
	17	M	C	D	D	D	D
	18	M	A	D	D	D	A
	19	M	E	C	C	D	D
Correct Questions				7 41	8 41	2 13	2 13

Table 2: Evaluation on Test 2

	Number	Difficulty	Correct Answer	Tensor w/ Reg	3-layer MemN2N
<i>Section 1</i>					
Long	13	M	A	C	train
	14	M	C	A	train
	15	M	E	D	train
	16	M	D	D	train
	17	M	D	B	train
	18	M	D	E	train
	19	M	E	E	train
	20	M	C	A	train
	21	M	A	B	E
	22	M	B	A	B
	23	H	B	B	B
	24	M	A	B	D
<i>Section 2</i>					
Medium	10	M	A	B	train
	11	M	A	B	train
	12	M	D	A	train
	13	E	C	C	train
	14	M	B	B	D
Medium	15	M	A	E	A
	16	M	E	A	train
	17	E	D	B	train
	18	E	A	A	train
	19	M	C	E	train
	20	H	E	B	train
	21	M	A	A	train
	22	M	C	C	E
	23	M	B	E	E
	24	M	B	B	E
<i>Section 3</i>					
Long	7	M	B	C	train
	8	E	B	E	train
	9	M	D	A	train
	10	M	D	C	train
	11	E	C	E	train
	12	M	A	E	train
	13	M	B	A	train
	14	H	E	E	train
	15	M	B	A	B
	16	M	E	E	E
	17	M	D	C	D
	18	M	B	A	B
	19	M	A	C	C
Correct Questions				11 40	7 14

Table 3: Evaluation on Test 3

	Number	Difficulty	Correct Answer	Tensor w/ Reg	3-layer MemN2N
<i>Section 1</i>					
Medium	10	E	B	C	train
	11	M	C	E	train
	12	E	A	B	train
	13	M	D	A	train
	14	M	C	A	C
Medium	15	E	D	C	C
	16	M	A	A	train
	17	H	B	A	train
	18	M	E	B	train
	19	M	D	A	train
	20	M	D	B	train
	21	M	D	C	train
	22	M	A	A	E
	23	E	E	E	E
	24	M	B	C	E
<i>Section 2</i>					
Long	13	M	E	A	train
	14	M	C	E	train
	15	M	D	A	train
	16	H	D	D	train
	17	M	B	A	train
	18	M	C	C	train
	19	H	B	B	train
	20	M	B	E	train
	21	H	D	B	B
	22	H	E	A	B
	23	M	A	A	E
	24	M	E	D	A
<i>Section 3</i>					
Long	7	E	D	B	train
	8	M	E	E	train
	9	M	D	B	train
	10	M	D	A	train
	11	M	B	A	train
	12	E	C	E	train
	13	M	A	A	train
	14	M	E	D	train
	15	M	C	A	E
	16	M	B	D	B
	17	M	E	C	D
	18	M	A	B	D
	19	M	D	E	D
Correct Questions				9	3
				41	14