Solving SAT Reading Comprehension Questions
with Memory Networks

Saahil Madge

Advisor: Professor Christiane Fellbaum

1. Background and Motivation

Machine question-answering (QA) has been a popular subject in recent research. It is a very interesting field from an academic perspective, requiring knowledge and techniques across the areas of artificial intelligence, machine learning, natural language processing, linguistics, and mathematics.

QA has applications to many aspects of daily life. A famous use of question-answering techniques is in Apple's Siri program. Siri combines speech recognition software with question-answering techniques to create a virtual iPhone assistant. Another example is IBM's Watson, which answered questions in many subjects and performed quite well as a contestant on the game show *Jeopardy!* A less spectacular but more commonly-used example can be seen by typing in a simple question such as *How many calories are in an apple?* on Google. The search results provide a detailed nutritional breakdown of an apple.

The general term QA covers a broad array of subfields. Information retrieval, factual question-answering, reading comprehension, and inference are all part of the greater question-answering area. QA is a rapidly evolving field, but as the examples above indicate, recent research has focused extensively on answering spoken questions, and factual questions.

However, the subfield of machine comprehension has not been studied nearly as well, though it has become quite popular in the past few years. Machine comprehension involves training programs and systems to read text and answer questions based on the newly acquired information. In most other QA tasks, the system has a large database of facts (knowledge

base) and must interpret the question and provide the relevant factual answer. Machine comprehension tasks, on the other hand, require the system to parse both the informational text as well as the question, and then provide the correct answer. As there is possibility for error within the knowledge base, machine comprehension tasks are some of the most difficult QA tasks.

Machine comprehension has a wide range of applications. Reading comprehension tests are a natural choice, but fields like document retrieval also use machine comprehension techniques. Financial firms that use news analysis to make trading decisions heavily rely on these techniques as well. Many long-term goals of AI such as dialogue and humanoid robots cannot be achieved without significant advance in this area [25]. Another incentive to study machine comprehension is that new research here can be applied to drive research in many other areas of artificial intelligence and natural language processing.

As a consequence of machine comprehension tasks being so difficult, most research with reading comprehension tests has focused on elementary school level tests. The most common dataset is MCTest [21], which we discuss in more detail in Section 3. This dataset is a set of passages and reading comprehension questions created by crowdsourcing. The stories are "carefully limited to those a young child would understand" [21]. Another common dataset is Facebook's bAbi dataset [25], which has a collection of short toy problems that can be used to train machine comprehension systems. These problems could certainly be solved by a human elementary school student. Berant et al. [3] train a system to answer questions based on a paragraph describing biological processes, which is certainly an advanced topic, but the system is specialized for biological tasks and not applicable as is to general reading comprehension.

In this paper we focus on SAT reading questions. The major reason is that they are far more complex than MCTest questions and other commonly-used machine comprehension datasets. The sentence structures, vocabulary, and topics covered in SAT passages are all far more varied than in the other benchmarks. The questions are much harder as well, and

simply understanding the question and the individual choices is difficult by itself. SAT questions often test underlying themes and broad ideas rather than particular factual details ("Why" or "How" questions). In contrast, MCTest questions are often of a "Who", "When", or "What" form, which have traditionally been easier for programs to answer.

Another aspect that makes SAT questions so difficult is that they require inference to answer, as opposed to just matching words or syntax. MCTest and similar questions can mostly be answered through word-matching (discussed in Section 2.1.1) or via syntactic matching. Word-matching means the system will try to find which sentences in the passage have the most words in common with the query, and use that to choose an answer. Syntactic matching is the same principle, but aims to match syntactic structure. SAT questions, however, almost never repeat the answer words in the question itself, and the query structure is not correlated with passage structure. To solve these questions, our system has to really understand the high-level concepts and information that is implied but not directly on the surface. This is a significantly harder problem than has been tackled before, but it is necessary to solve to continue moving forwad in the field of machine comprehension. Essentially, by trying to answer SAT questions we can try to solve problems that will be present in real-world applications but have not yet been tackled by existing research.

Due to the difficulty of interpreting the text as-is, we represent the text using relational logic instead. Essentially we dynamically create a *Knowledge Graph* (covered in more detail in Section 2.1) as we parse through the text and try to frame each question as a query on this knowledge graph. The knowledge graph representation is typically used for factual data that has already been collected in entity-relation triple form. As far as we know, no previous research in this space has tried to dynamically create a knowledge graph on which we can answer questions.

Traditionally most research on relational learning has used either neural networks or tensor decomposition to answer queries on the graph. In this work we present an improved version of both methods. For tensor factorization we enhance a method known as RESCAL

[1, 18] with a technique known as *Semantically Smooth Embedding* [10]. RESCAL trains a factorization of the knowledge graph using Alternating Least Squares, and this factorization can then be used to answer queries.

The neural network approaches such as TransE [6] and TransH [24] use neural networks to train an embedding method for the relations on the knowledge graph, and use the trained model to answer queries. However, none of these previous approaches are designed for question-answering. We enhance the Memory Networks architecture proposed by Weston et al. [26] and improved by Sukhbaatar et al. [23] to take knowledge graph triples and queries as input, and provide a scoring of the answer choices as output. We discuss memory networks and alternatives in much greater depth in Section 3 and Section 4.

Our goal with this research is to present a way to combine ideas that were previously used in isolation for QA or other learning tasks into a system which can parse and interpret a complex text, and answer difficult high-level questions about the text. The text, questions, and techniques we use are more ambitious than those studied in previous research. Our results should be interpreted as merely proof-of-concept. We hope that this paper will encourage others to tackle problems of equal or greater ambition, and drive progress in the development of safe artificial intelligence.

2. Technical Review

In this section we review some of the techniques that are used in previous research and in this paper.

2.1. Representation Techniques

We begin by reviewing various methods for embedding text and performing queries on the embedded representations.

2.1.1. Bag-of-words

Bag-of-words is a type of analysis in which we treat each sentence as simply a collection of words, without paying attention to their order or the dependencies between individual words. If we have some vocabulary of size V, we can encode each sentence as a vector of size V where each vector element corresponds to some word in the vocabulary. If the word appears in the sentence we store how many times it appears, and if it does not appear we store it as a 0. This is known as *one-hot* vector encoding. Most designs encode each input sentence and the query as one-hot vectors. These vectors are not used as-is, but are first embedded into some dimension d. The learning algorithm is run on the embedded vectors, and we can output a response in a specified format. Two common methods are to return a probability vector of size V where each element represents how likely that word is to be the output, or to pick the most likely word from that distribution and return only that word.

2.1.2. Dependency Parsing

Dependency parsing is an important technique used to understand the structural information in a sentence. Essentially, dependency parsing creates a graph of the words in the sentence and describes how each word relates to other words, using the standard Universal Dependencies [9]. A sample graph is shown in Figure 1, from [8]. Each dependency in the graph has a *governor* and *dependent*. The collection of dependencies contains all the structural relationships in the sentence. Each dependency is labelled to show the exact relationship between governor and dependent. For example, in the graph in Figure 1 the label *nsubj* highlights the nominal subject of a particular clause in the sentence [9]. The specific parser we use is the Stanford High-Performance Dependency Parser by Chen and Manning [8].

Once the dependency graph is created, we can analyze it to extract the relevant information in the sentence. The exact details are discussed more in Section 4.1.2. The basic approach is to find the main subject of the sentence and either the direct object it acts on or some

Figure 1: A small example knowledge graph, from [8]

description of it. Then it can be turned into an entity-relation triple on the knowledge graph.

2.1.3. word2vec

The *word2vec* embedding technique introduced by Mikolov et al. [16] is perhaps the most common embedding technique used in NLP research. It is actually an umbrella term for various methods, such as *skip-gram* or even bag-of-words embedding. The word2vec techinque produces a vector representation of words, often with several hundred elements. These representations can then be manipulated as vectors in the higher-dimensional space, allowing for interesting operations on words. For example, **vector("king")** - **vector("man")** + **vector("woman")** produces a vector that is very close to **vector("queen")**. The reason word2vec is able to embed so successfully is generally thought to be the fact that the process of embedding gives similar words a similar embedding (which is measured by cosine similarity). The cosine similarity measure is also used to find the similarity between two arbitrary words. *word2vec.similarity*(word1, word2) is a function within word2vec which returns a measure between 0 (not at all similar) and 1 (identical).

2.1.4. Knowledge Graph Overview

A knowledge graph generally refers to a set of *entity-relation triples*. These triples take the form $\langle e_1, r, e_2 \rangle$. where e_1 and e_2 are entities and r describes the relation between these two entities. Entities and relations (collectively referred to as *labels*) can be repeated across triples. The total set of triples denotes all the knowledge in our "Knowledge Base".

We can represent these triples as a directed graph, with each entity as a vertex and the

Figure 2: A small example knowledge graph

relations as edge types between them. Figure 2 shows an example of a very small knowledge graph. Note that there is only one node per entity. This allows the graph to store multiple pieces of information regarding the same entity together, to allow for easy processing by the program while also making it simple for humans to understand. The mathematical details of our knowledge graph model can be found in Section 4.

Once a knowledge graph is built, we can pose queries about the graph. For example, we can provide two entities e_1, e_2 and ask which relation r links the two on the graph. We can also provide an entity e and the relation r, and ask for which other entity e' do the triples $\langle e, r, e' \rangle$ or $\langle e', r, e \rangle$ exist on the graph. Put another way, the queries are either *link prediction* queries or *entity prediction* queries.

2.2. Neural Networks

As discussed in Section 1, recent research in machine comprehension has extensively used neural nets (NNs) and their variants. Here we provide a quick review of neural network

design. These topics can be covered in far more depth in [5, 19].

2.2.1. Feed-forward Neural Networks

The most basic model is known as a *feed-forward* neural net. These have several layers composed of units (nodes). There is an initial input layer, a final output layer, and *hidden layers* in the middle. A unit in layer l takes as input the output of nodes in the previous layer (or the actual input if this is the input layer). It multiplies these by some weight matrix W^l, transforms it using an activation function g, and sends its output to the next layer (or as the final output if this is the output layer). Common activation functions are the *tanh* and *sigmoid* functions. The sigmoid function is defined as $\sigma(x) = \frac{1}{1+e^{-x}}$. For now we will assume that our activation function is sigmoid.

Mathematically, we define the input to the current layer as x^{l-1}, where x_{ij}^{l-1} means that it is the input from unit j in the previous layer to unit i in the current layer. W_{ij}^l is the weight on this connection, defined as 0 if node i does not depend on node j. The output of node j is $\sigma(W^l x_j^{l-1})$. Typically the quantity $W x_j^{l-1}$ is known as z_j. We can also define bias vectors b such that $z_j = W^l x_j^{l-1} + b_j^l$. To summarize, we can write the output of some layer l in vector notation as

$$o^l = \sigma(W^l x^{l-1} + b^l)$$

We can propagate this forward starting at the input layer, through the hidden layers, and finally at our output layer. We can then train our network using the backpropagation algorithm. We have some cost function C that is a measure of the error of our output. We find the error at the output layer, and then change the weights at each previous layer, working backwards to the input layer. The key to backpropagation relies on updating the weights based on their partial derivatives with relation to C. In-depth derivations are provided in [5, 19].

2.2.2. Recurrent Neural Networks

Feed-forward networks are good tools, but do not have any form of state (memory). Here we review recurrent neural networks (RNNs) which have a hidden state h_t.

At step t in a recurrent neural network, we have input x_t and a hidden state h_{t-1}. We also have weight matrices U and W. We can define the current state $h_t = \sigma(Ux_t + Wh_{t-1})$. We have another weight matrix V for the output. The output $y_t = softmax(h_t)$, where $softmax(z_i) = \frac{e_{z_i}}{\sum_j e^{z_j}}$. The softmax function output is a probability distribution over the vector.

Using the softmax function has several advantages. Because it is a smooth function, we can take its derivative, making it easier to calculate the backpropagation equations. Additionally, most RNNs and variants use a loss function known as cross-entropy loss. If the model output is $\hat{y}$ and the real (target) is y, then cross-entropy loss is defined as $\sum_i y_i \log(\hat{y}_i)$. When we pair a softmax output with cross-entropy loss, the error at the output node i is simply $\hat{y}_i - y_i$. This is very convenient because it can be calculated rapidly, so it allows for fast training of the RNN. Training is normally a bottleneck, and by combining a robust output function and robust loss function we get a very elegant output error. As this is a very important result, we have included it in Section 8.1.

3. Related Work

The first major research in machine comprehension was conducted by Hirschman, et al. [11]. Their system "Deep Read" takes a story and a series of questions as input, and proposes answers by using bag-of-words techniques with other methods such as stemming and pronoun resolution layered on top. On a collection of remedial reading questions for grades 3-6, Deep Read answered about 40% of the questions correctly.

Most recent research has used the MCTest [21], released by Richardson, et al. at Microsoft research. There are a total of 500 passages, each with 400 questions, created at the reading level of a young child. The stories are fictitious, which means the answers can only be found

in the story itself. Richardson et al. also provide two baseline implementations to serve as a benchmark. The first is a simple bag-of-words model with a sliding window. It scores about 51% accuracy on the MC500 questions. The second implementation adds a distance-based scoring metric and reaches 56% accuracy. Both implementations score significantly higher on questions where the answer is contained in just one sentence than on questions where the answer requires information across multiple sentences.

Narasimhan and Barzilay [17] also work on the MCTest tasks. Their main insight is to create a task-specific discourse parser to capture specific discourse relations in the MCTest passages, while the prevailing method had been to use generalized off-the-shelf discourse parsers. They create three models of increasing complexity. The first assumes each question can be answered using just one sentence, and estimates a joint probability distribution across the question, query, and answer. The second model adds in the joint distribution across a second sentence as well, to handle the case in which the answer needs two sentences to answer. The third model incorporates discourse relations between the two sentences to better understand the relationship between information in each sentence. The relations are defined as "Causal", "Temporal", "Explanation", and "Other". Most systems have performed the worst on these types of questions. By focusing specifically on modeling these explanatory relations, the third model easily outscores the other two, and the two MCTest baseline implementations, achieving almost 64% accuracy.

Berant et al. [3] also focus on analyzing inter-sentence relations, with application specifically to passages and questions about biological processes. These passages generally describe a chemical reaction or other process in which there are various starting entities which interact with each other and form a new output. Understanding how the inputs interact and tracing the flow of the process is crucial to answering the question. To solve this, Berant et al. define events (or non-events) as "triggers", and try to find relationships between events. The events can be thought of as nodes on a graph, with an edge defining some relation between the two. There are eight possible relations, including "cause", "enable",

10

and "prevent". They first create events and then predict relations between the events. The queries are also formulated as a graph. They are categorized as dependency questions, temporal questions, or true-false questions. This model scores almost 67% on the dataset, over 6% better than the next-best model.

Most recent machine comprehension research has focused on using neural nets as the core system. Neural nets are a very generalizable technique and in the past decade have become very popular. In fact, advances in neural net techniques have greatly contributed to the recent surge in machine comprehension research.

Weston et al. [26] introduced memory networks. These are a type of neural network which simulate a long-term memory. We discussed in Section 2.2.2 how RNNs have a "vanishing gradient" problem, and are not able to take advantage of states from more than a few steps prior. Memory networks have a memory module which stores old memories and then updates them given new inputs. When given a query, the memory network finds relevant memories, then finds memories that are relevant given those memories, and so on. Finally, it provides an answer to the query. Sukhbaatar et al. [23] improved this model by creating a model that can be trained end to end, while in the original design each module needed to be trained independently. This is the model that is the basis for our design, so we discuss it in more detail in Section 4. Kumar et al. [14] create "Dynamic Memory Networks". Their design uses two memory modules. The semantic memory module stores general knowledge, while the episodic memory module iteratively finds memories relevant to the query.

Knowledge Graphs are also a popular research area. There are two main approaches that are used for solving knowledge graph embedding and query problems. The first is the neural network embedding style, popularized by Bordes et al. [6] and improved by Wang et al. [24]. Bordes' TransE and Wang's TransH both train a neural network to recognize triple embeddings in some higher dimension. The triples are embedded in such a way that triples that are "similar" according to some metric are embedded near each other. Queries can be

answered more easily using the embedded vectors.

The second approach focuses on representing the knowledge graph as a tensor. Nickel et al. [18] proposed RESCAL, a model which converts the knowledge graph into a tensor of adjacency matrices and trains a latent rank-r factorization of the tensor using a technique called ASALSAN [1]. Queries about the knowledge graph can be easily converted into searching over a relevant subset of elements in the factorization. Chang et al. [7] created TRESCAL, which is an optimized version of RESCAL.

4. Model

This research presents three fundamental insights.

1. We treat the input text as a knowledge base and dynamically convert it into a knowledge graph. We can then interpret the questions as queries on the graph, requiring us to find an entity or relation that represents the answer to the question.

2. We use tensor decomposition to embed the knowledge graph and questions. Specifically, we use an enhanced version of RESCAL [18]. RESCAL by itself is a fine model, but we make it more accurate by adding "Semantically Smooth Embedding"(SSE), as proposed by Guo et al. [10]. In their original paper they propose adding this type of embedding to tensor decomposition as future work, so to our knowledge we are the first to perform tensor decomposition with SSE.

3. We embed the knowledge graph and questions into input vectors, and pass them into a Memory Network. Given a set of embedded inputs and queries, this type of neural network takes advantage of "memory" to find new facts that are relevant given the current set of relevant facts, but may not be relevant solely based on the query.

These techniques were first introduced in previous research. However, all of them were applied to separate domains. Neural nets have been used for machine comprehension on MCTest, bAbi, and questions of similar difficulty. Their embedding model is quite simple, however (often just bag-of-words). Knowledge graphs and tensor decomposition have traditionally been used only for relational learning when a knowledge graph or set of triples already exists. As our problem is more complex than that tackled by these original models, and we must combine these original models, we have to significantly modify their conceptual and mathematical basis.

4.1. Knowledge Graph Representation

The first part of our design is to dynamically convert the input passage into a knowledge graph and the questions as queries on that graph.

4.1.1. Advantages of Knowledge Graph Representation

There are several advantages to using a knowledge graph representation. The main benefit is that it actually preserves the meaning of the text. Whereas bag-of-words loses the semantic and structural information in a sentence, a knowledge graph representation almost completely stores the intended meaning of the original text. Most sentences in English can be broken down into Subject-Verb-Object triples, which require little if any processing to convert into the $\langle entity_1, relation, entity_2 \rangle$ format. The knowledge graph itself is simply a list of these triples. If we can parse the input text and extract the Subject-Verb-Object relationship, we can convert each sentence into a triple of the form

$$\langle subject, verb, object \rangle \Rightarrow \langle e_1, r, e_2 \rangle$$

The details of the conversion are covered in Section 4.1.2.

Even more importantly, the knowledge graph inherently has *memory*. What this means

is that it automatically pieces together information about the same concept, even if that information occurs across different sentences or even different paragraphs. Consider a passage about a boy named Harry. Even if in one paragraph the text says that Harry was born in 1980 and several paragraphs later we learn that he has a friend named Ron, these relations both link to the single Harry entity node. By analyzing the outbound and inbound relations of the Harry entity node, we can easily find all the information the text has stated about Harry. Any question asked about Harry can be answered using this information. While such a representation for memory may seem natural, virtually all previous research in the machine comprehension space has parsed by sentence. As a result, the knowledge graph representation is the best way for our system to easily find all information about an entity that was given in the original text.

Entities also do not have to be concrete things; they can just as easily be abstract concepts. For example, Figure 2 shows entities like Carl Pavano (a person), but it also shows an entity node for the year 1976. To roughly generalize, any noun can be an entity node. The phrase "abstract concept" itself could be an entity node. Entity-relation triples can all be analyzed similarly, so abstract reasoning is as simple as concrete reasoning. Not only does a knowledge graph allow us to perform abstract reasoning, but it is arguably the only way to perform reasoning of this form.

A good way to understand the knowledge graph representation is to think of it as a *concept map*. Every sentence in the text introduces a new concept (concrete or abstract) and/or provides more information about a previously introduced concept. All information is provided in such a way that the concept is either the actor or the receiver of some action. We update our concept map with this new information by adding entity nodes as needed and the relation provided in the current sentence. When we've finished parsing the input text, we have the full concept map, which gives complete information on each concept and the information we know about it.

Another benefit is that there are already methods and theory of analysis for knowledge

graphs. As mentioned earlier, passages as complex as SAT reading comprehension tests have rarely been analyzed semantically. Moreover, the fact that information comes in a variety of sentence structures across several paragraphs means that traditional semantic analysis will not suffice. Traditional analysis often parses one sentence or one paragraph at a time, which is not enough for our particular use case. The original input simply cannot be analyzed as it is using existing techniques. Since this problem cannot be solved in its original form, we convert it to a form which we know how to analyze. Not only have knowledge graphs been studied extensively in the field of relational learning, but they can also be analyzed using techniques from graph theory. Essentially, we turn the problem of analyzing complex text into a problem of analyzing a knowledge graph.

At the same time, the knowledge graph representation also preserves human readability. A major problem with text input is that it is easily interpretable by humans but requires much preprocessing before it can be interpreted by a program. It is not necessary for our processed text to be in a format humans can interpret, but it certainly helps us reason about the problem and come up with interesting solutions. The entity-relation triple format is almost as intuitive to humans as the original textual input.

Of course, even the knowledge graph representation has some drawbacks. Because we build the knowledge graph using the statements in the original text, we can only perform analysis on the textual information. We can trace relations through the graph but there is no way to perform higher-level reasoning. For example, questions that ask about classification of new information or broad themes of the passage cannot be answered. Nevertheless, these are flaws in any form of representation. The knowledge graph representation provides the most advantages and fewest drawbacks of any representation.
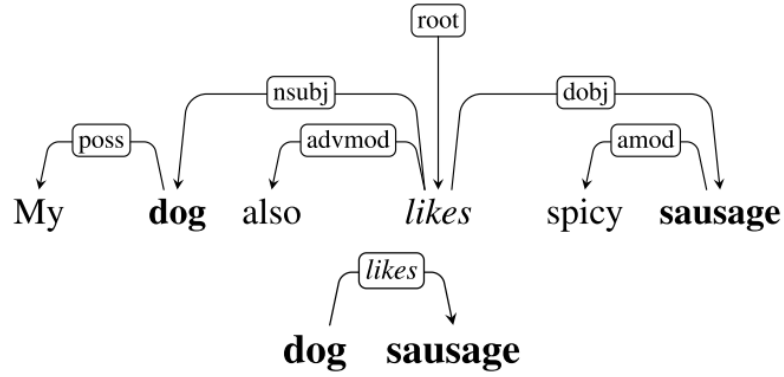
Figure 3: Extracting an entity-relation triple from a sentence, from [17]

4.1.2. Converting to Knowledge Graph Representation

In this section we present the exact details of converting the text into a Knowledge Graph. Section 5 discusses how the raw tests are converted into textual information that we can then parse.

We use the Stanford High Performance Dependency Parser [8] to extract the entity-relation information from the text. Section 2.1.2 explains how this tool creates a dependency graph of a particular sentence. We then provide a set of semantic rules to isolate the important information and create a triple from it. Figure 3 from [17] shows a simple sentence, as well as the extracted triple.

The information extracted should be in knowledge graph format, namely $\langle \text{entity}_1, \text{relation}, \text{entity}_2 \rangle$. We want to frame the information in the sentence so that it is one entity with some relation to another entity. Most sentences are of the form where the subject performs some action or has some description, so we focus on these two types of sentences. In each sentence we isolate the main subject using the *nsubj* dependency. Then we check for one of two relations. If the subject is performing some action then the sentence can be generally extracted as $\langle \text{subject, action, direct object} \rangle$. In this case the *dobj* dependency is used to extract the action verb as well as the direct object. Figure 3 shows this case. In the second case, the sentence is often of the form "<subject> is... <description>". Here we use the *cop* dependency to extract the verb and the description. The triple $\langle \text{subject, copular verb, description} \rangle$ is added

16

to the knowledge graph. In English, "verbs of being" such as "is", "are", "were", etc. are all copular verbs. Narasimhan and Barzilay [17] do a similar type of extraction, but they create what they call an *entity graph*, which is slightly different from our knowledge graph. Their entity graph refers to linguistic entities from the dependency parsing, while our knowledge graph refers to information, not linguistic objects.

4.2. Tensor Decomposition

*******REVISE THIS PART AND ALSO DESCRIBE HOW THE WORD2VEC QUESTION MATCHING IS DONE TO CONVERT QUERIES*******

First we discuss the tensor decomposition approach to answering SAT questions. The advantage of using tensor decomposition is that it is based on well-understood mathematics, rather than simply empirical validation as with the neural network.

Our model is an improved version of RESCAL, proposed by Nickel et al. [18] and used by Chang et al. [7]. RESCAL itself takes the general tensor decomposition model proposed by Bader et al. [1] and modifies it to work for relational learning.

4.2.1. Knowledge Base Tensor

The tensor decomposition approach treats the knowledge graph as a directed graph. A tensor is simply a multi-dimensional array. In our case, we use a three-dimensional array $\mathcal{X}$ to represent the knowledge graph. If we have N entities $e_1, e_2, ..., e_N$ and M relations $k_1, k_2, ..., k_M$ then we can consider $\mathcal{X}$ as M layers of $N \times N$ matrices. Each layer $\mathcal{X}_k$ tells which entities are linked by relation k. If we denote our knowledge base as KB, then $\mathcal{X}$ is formulated as

$$
\mathcal{X}_{ijk} = \begin{cases} 1 & \text{if } \langle e_i, k, e_j \rangle \in KB \\ 0 & \text{otherwise} \end{cases}
$$

Figure 4 from [7] presents a graphical view of $\mathcal{X}$. $\mathcal{X}$ is essentially a series of adjacency

17

Figure 4: A graphical representation of $\mathcal{X}$, from [7]

matrices $\mathcal{X}_k$, where each matrix only marks edges of relation type k.

4.2.2. Tensor Factorization Overview

Our approach aims to factorize the tensor and take advantage of the inherent structure of the graph. To that end we train a rank-*r* factorization of each slice, as suggested in [1, 18].

$$X_k \approx AR_kA^T \text{ for } k \in [1,m]$$

Here, A is an $n \times r$ matrix which contains the latent-component representation of the entities in the knowledge graph. We can think of this as an *r*-dimensional embedding of the entities, where each entity has *r* hidden features. This is conceptually similar to the word2vec embedding, where each feature of a word vector does not necessarily mean something intuitive but is a hidden component found through training.

R_k is an $r \times r$ matrix which represents how these hidden features interact with each other for the *k*-th relation. R_k is asymmetric, which means that the hidden feature interactions are

18

one-way. It is not necessarily true that $R_{ijk} = R_{jik}$, although it could be. Note that R itself is also a three-dimensional tensor of shape $r \times r \times k$.

We decompose the tensor so that we can answer queries using the factorization, which has the hidden components of the graph and is thus more generalizable, as opposed to using the original tensor which is simply an adjacency tensor of the knowledge graph. A particular connection between two entities, for example, might not be explicit on the knowledge graph but can be uncovered by the factorization.

We frame the problem of finding A and R_k as an optimization problem, where we aim to minimize the difference between the original matrix $\mathcal{X}_k$ and its factorization AR_kA^T. Specifically, define the loss function f as the squared frobenius norm of this difference.

$$f(A, R_k) = \frac{1}{2} \sum_k \|\mathcal{X}_k - AR_kA^T\|_F^2 \tag{1}$$

where the frobenius norm of some $a \times b$ matrix X is defined as

$$\|X\|_F = \sqrt{\sum_{i=1}^{a} \sum_{j=1}^{b} |X_{ij}|^2}$$

To find the optimal A and R_k, we just have to minimize 1, along with some regularization function $g(A, R_k)$. Formally,

$$\min_{A, R_k} f(A, R_k) + \lambda g(A, R_k) \tag{2}$$

In [18][7] the regularization term is simply defined as $g(A, R_k) = \frac{1}{2}\lambda(\|A\|_F + \sum_k \|R_k\|_F)$. This is a standard regularization function where λ is a hyperparameter. The aim is to prevent overfitting when we try to solve the optimization problem. We enhance their approach by using a stronger regularization function which adds more constraints on the factorization.

4.2.3. Semantically Smooth Embedding

Semantically Smooth Embedding (SSE) was proposed by Guo et al. [10] as a way to improve the TransE [6] embedding model. Whereas the traditional models embed each individual fact in isolation, SSE imposes geometric constraints on the whole model.

Specifically, we modify the *Locally Linear Embedding* (LLE) constraint. This constraint states that each node should be roughly a linear combination of its neighbors [22]. Guo et al. [10] created a version meant for neural net embeddings such as TransE, but is not applicable to the tensor decomposition method. Here we review their version first and then present our new version for tensor decomposition.

In the original implementation each entity was labeled with a category and then embedded into a vector representation. All nodes in the same category as e_i were defined as its neighbors and this set was denoted $\mathcal{N}(e_i)$. Then they defined

$$
w_{ij} = \begin{cases} 1 & \text{if } e_j \in \mathcal{N}(e_i) \\ 0 & \text{otherwise} \end{cases}
$$

$$
R = \sum_{i=1}^{n} \| \mathbf{e_i} - \sum_{e_j \in \mathcal{N}(e_i)} w_{ij} \mathbf{e_j} \|_2^2
$$

where R is the regularization term added to the loss function. The logic here is that we suffer more penalty in our loss function when entities that are close to each other in the original set (i.e. are neighbors) have embedded vectors that are far from each other. We cannot use this equation for our purposes since our entities do not have categories and we do not have embedded vectors, but we can apply similar logic.

We propose defining the neighbor set of each entity e_i in our graph as all other entities for which there is a directed edge from e_i to e_j in the knowledge graph. The full neighbor set

matrix is denoted W.

$$w_{ij} = \begin{cases} 1, & \sum_k \mathcal{X}_{ijk} \geq 1 \\ 0, & \text{otherwise} \end{cases} \tag{3}$$

$$e_j \in \mathcal{N}(e_i) \text{ iff } w_{ij} = 1 \tag{4}$$

That is to say, $w_{ij} = 1$ if and only if we have some relation r for which the triple $\langle e_i, r, e_j \rangle$ exists in the knowledge base. We do not consider entities which are related to e_i in such a way that e_i is the recipient of the relation. Note that this means the neighbor relation goes one way. $e_j \in \mathcal{N}(e_i)$ does not imply that $e_i \in \mathcal{N}(e_j)$.

Now consider the nodes in $\mathcal{N}(e_i)$. Since e_i is related to all of these, and because we want to enforce some geometric structure on our representation (as done by Guo) we claim that the latent properties of each node e_i, as provided in the matrix A, should be some linear combination of its neighbor set.

$$A_i \approx \sum_{i=1}^{n} w_{ij} A_j = \sum_{i=1}^{n} W_i A$$

We can penalize more for an A in which this distance is large. The reason we made the neighbor relation only one way is that we do not expect a large number of disconnected components of our knowledge graph, and a two-way neighbor set would push every node to have the same latent representation.

4.2.4. Solving the Optimization Problem

The new, semantically smooth regularization function is

$$g(A, R_k) = \sum_{i=1}^{n} \left\| A_i - \sum_{j=1}^{n} w_{ij} A_j \right\|_F^2 = \sum_{i=1}^{n} \| A_i - W_i A \|_F^2 \tag{5}$$

We are now ready to actually solve for A and R_k. We want to minimize $f(A, R_k) + g(A, R_k)$. This is a nonlinear, non-convex optimization problem which would require complicated techniques. Calculating f and g is expensive, so we use the Alternating Least Squares(ALS) method [13]. Specifically, we use the ASALSAN factorization method [1].

ALS turns the non-convex optimization problem into a convex problem by fixing one of the unknowns. Then the problem can be solved with ordinary least-squares optimization, so we use the normal method of least-squares to update the unknown variable. Then we alternate by fixing this unknown and changing the other variable. The ordinary least-squares can be solved relatively fast (although calculating inverse matrices is expensive) and since each unknown is calculated independently (keeping the other fixed), the ALS algorithm can be easily parallelized. These two aspects make it very computationally fast, a quality necessary for our purposes.

We have

$$\bar{\mathcal{X}} = A\bar{R}(\mathbf{I}_{2m} \otimes A^T)$$

where $\otimes$ is the Kronecker product and

$$\bar{\mathcal{X}} = \left(\mathcal{X}_1 \mathcal{X}_1^T ... \mathcal{X}_m \mathcal{X}_m^T \right)$$

$$\bar{R} = \left(R_1 R_1^T ... R_m R_m^T \right)$$

Then by ASALSAN we derive the following update rules. We first update A while keeping R_k fixed.

$$A \leftarrow \left[\sum_k \mathcal{X}_k A R_k^T + \mathcal{X}_k^T A R_k \right] \left[B_k + C_k + \mathbf{I} \right]^{-1}$$

where

$$B_k = R_k A^T A R_k^T, \quad C_k = R_k^T A^T A R_k$$

22

To update R_k, we first fix A. Then we notice that if we vectorize $\mathcal{X}$ and R_k, we can rewrite Equation 2 as

$$f(R_k) = \|\textbf{vec}\,(\mathcal{X}_k) - (A \otimes A)\,\textbf{vec}\,(R_k)\|$$

where we have removed an extra term $\lambda\,\|\textbf{vec}(R_k)\|$ because R_k is no longer in our regularization term. This rewritten f is actually a linear regression problem whose solution is

$$R_k \leftarrow \left((A \otimes A)^T (A \otimes A) + \mathbf{I}\right)^{-1} (A \otimes A)\,\textbf{vec}\,(\mathcal{X}_k)$$

We alternate updating A and R_k until we have performed some maximum number of iterations or until $\frac{f(A,R_k)}{\|\mathcal{X}_k\|_F^2}$ is below some tolerance ε. The final factorization $\hat{\mathcal{X}}$ can now be used to answer queries.

4.2.5. Answering Queries

The first step in answering SAT questions with tensor decomposition is to convert them to queries on the knowledge graph. Queries are either link prediction queries or entity prediction queries (Section 2.1.4). We use *word2vec* to perform the conversion. As each entity and relation is a word, we test the similarity of every word in the question to every entity and relation. The two highest scoring graph labels (at least one entity) are the knowns in our query. If there are two entities, then we have a link prediction query. If there is an entity and a relation, then we have an entity prediction query. The advantage of using word2vec similarity is that SAT questions often rephrase the textual words. For example, if in the passage a character is "angry", the question will use a word such as "incensed", which does not exist in the knowledge graph, instead. word2vec translates unseen words into knowledge graph labels which can then be operated on.

Link-prediction queries are relatively simple to answer. Given e_i, e_j we return the relation

k that satisfies $\underset{k}{\arg\max}\, AR_k A_{ijk}^T$. This is the relation which our factorization guesses is strongest between the two entities.

Entity prediction queries are answered similarly. Given e_i, k we return the entity e_j that satisfies $\underset{j}{\arg\max}\, AR_k A_{ijk}^T$. This is the entity which our factorization guesses is most likely to be the recipient of relation k from entity e_i.

Once we find the label that satisfies the query, we test the similarity of this label to every word in every answer choice, and pick the answer choice that matches best. To summarize, we first convert the question a query on the knowledge graph, find the label in the knowledge graph that answers the query, and choose the answer that is most similar to it.

4.3. Memory Networks

The second way we propose answering queries is with neural networks. Specifically, an enhanced version of Memory Networks. The advantage of the neural network approach is that it is very intuitive, computationally feasible, and empirically proven. However, it does lack the strong theory of the mathematical approach. We choose the Memory Networks architecture as the basis for our model because it has strong performance on the bAbi dataset and the foundational architecture itself scales easily to more complex problems as long as the input is in a valid format. First we review the architecture as proposed in [23]. Then we discuss how to adapt it to this specific problem.

4.3.1. Memory Network Architecture

First we will review the architecture for one layer, as the extension to several layers is straightforward. The basic model takes in a set of input vectors $x_1, ..., x_n$ and a query vector q, and outputs an answer a. Memory Networks require the input text and the query to be vectorized in some form for easy computation. These embedding techniques are covered in detail in Section 2.1. Typically the embedding is done with bag-of-words embedding. This embedding technique relies on the same word being used in the answer choices as well

as in the original text, which is true for MCTest and bAbi. The vocabulary comes from a dictionary of fixed size V.

The input vectors are converted into a set of memory vectors m_i using a matrix A of size $d \times V$, where d is a pre-determined dimension. q is also embedded into a state u using a different matrix B of size $d \times V$. To find which memory vectors are most relevant to the initial query, we take the dot product of each vector with u.

$$p_i = \text{Softmax}(u^T m_i)$$

where $\text{Softmax}(z_i) = \frac{e^{z_i}}{\sum_j e^{z_j}}$. Since the softmax function returns a value between 0 and 1, each score p_i is a scalar that can be considered a probability of how relevant each memory vector is to the query. The more relevant it is, the more likely it will help find the answer.

The input vectors are now sent through another matrix C to produce a set of transformed inputs c_i. Each transformed vector c_i is weighted by its probability c_i and the sum is taken as the output o of that layer.

$$o = \sum_i p_i c_i$$

This output can then be fed as an input to the next layer, but in the single layer case it is added to u and sent through one final matrix W of size $V \times d$. Lastly, a softmax is taken on the result to yield a probability distribution over the vocabulary space. This is designed to pick one word or set of words as the answer.

$$\hat{a} = \text{Softmax}\left(W\left(o + u\right)\right)$$

The network itself is trained using backpropagation. The answers to the training set questions are provided as input, and the network learns to answer them. The exact details of training are covered in Section 6.2.

Several strategies can be used to extend the Memory Network to multiple layers. We

use the *Recurrent Neural Network* strategy as it makes most sense for our purposes. This strategy is the simplest to understand as well as to train. Each layer acts the same as in the one-layer case. The only difference is that for all layers $k > 1$ the input $u^{k+1} = u^k + o^k$, and the matrix W is only applied to o^{final}. Each layer uses the same A, B, C matrices. By using the same matrices at each level and passing the state along, this architecture acts like a Recurrent Neural Network.

The initial layer calculates how relevant each memory vector is to the initial query. That output is combined with the initial query and passed as u to the next layer. That layer finds how relevant each memory vector is to the new state u. This means that memories which were not very relevant to the initial query may be relevant to the memories which were relevant to the initial query. Each layer finds new memories that can help answer the question.

4.3.2. Memory Network Enhancements

As discussed, the original Memory Networks model uses simple bag-of-words embedding to vectorize the text into the input vectors x_i. The problem with this method is that the "most relevant" memories found in each layer are relevant based on the exact word choice, not based on conceptual similarity. To answer SAT questions the text must be vectorized in such a way that the knowledge graph relationships are captured, and the question must be converted into a query vector for the network. We need to change the inputs while preserving the overall architecture so that the memories chosen as relevant actually match the concept in the query, as opposed to simply matching the words.

We propose stacking horizontally each slice of the tensor $\mathcal{X}$ as defined in Section 4.2.1 and using each column as an input vector. In total the input matrix $\mathcal{X}'$ will be of shape $n \times (nm)$. The query is more complicated to vectorize. The memory network architecture requires the query to be a 1-dimensional vector. We create a vector q of length $N + M$ where N is the number of entities and M the number of relations, where each element is 0. We find

the two most similar labels as in Section 4.2.5 and set their corresponding vector elements to 1. This is a two-hot encoding vectorization.

We make two changes to the original network architecture. We set the current layer $u^k = R(u^{k-1} + o^{k-1})$, and o is now the dot product of Cx_i and the p_i vector, instead of being a scaled sum. Adding another matrix R allows us to capture extra complexity across layers while preserving dimensionality. Changing o to be a dot product is also to help with dimensionality. The original matrix had a fixed query size as well as input length of 177, but with SAT questions the length $N + R$ is different for every knowledge graph input, so we need to add more flexibility in the architecture for this changing dimensionality.

The architecture equations become

$$m_i = Ax_i$$

$$u^0 = Bq, u^{k+1} = R(u^k + o^k)$$

$$p_i = \text{Softmax}\left(\left(u^k\right)^T m_i\right)$$

$$c_i = Cx_i$$

$$o^k = \sum_i p_i c_i$$

$$\hat{a} = \text{Softmax}\left(W\left(o^l + u^l\right)\right)$$

where l is the total number of layers.

In the multiple-layer case, the approach is the same as in the original model where the matrix C stays constant at each layer, and $u^{k+1} = u^k + o^k$. Specifically we use a 10-layer neural network. The original paper used a 3-layer model to find 3 steps of relevant facts. SAT questions require information pieced together from many sentences, however, so we add more layers to capture this additional complexity.

27

4.4. Model Comparison

The Tensor model and the Memory Network model have the same underlying goal but go about it very differently. Both models try to find the inherent properties of the knowledge graph given the graph itself. Essentially, by finding this latent representation of the graph, both models can answer questions about it.

The tensor model has a much stronger mathematical base. The RESCAL model uses ASALSAN [1] which is a general mathematical model for tensor evaluation. Using ALS to optimize is an approximation to make the problem computationally feasible, but the technique is still mathematically sound. At the same time however, the tensor model is trained only on the knowledge graph itself and not on the answers. IT is heavily reliant on the thoroughness of the knowledge graph triples extracted from the text. Information that is missed in the triple extraction will be lost to the tensor model as well.

The Memory Networks model is the exact opposite. It has a much weaker theoretical base but is also more robust to errors in the knowledge graph extraction process. The mathematics of neural networks are not as well understood as that of tensors, especially as we add more layers. This means that the neural net model we are using can behave in various unexpected ways. Nevertheless, the model is trained using the answers to the training set questions. It can find patterns and information using the training process, and as a result is less subject to fluctuations in the actual knowledge graph extraction.

5. Data

The model evaluation is conducted on SAT tests in the old (pre-2016) format. There are several reasons we use the old format. The main reason is that the old reading section consists of mostly reading comprehension, while in the new format there are charts and graphs that must be interpreted as well. Reading these requires some computer vision work on top of our existing models, and the questions do not naturally make sense as queries on a knowledge graph. The old format was tried and tested for many years, but the new format is

just in its first year. Additionally, there are far more resources for the old SAT which we can use for data.

There are no publicly available datasets for the SAT. We tried contacting ETS to obtain a research dataset, but as we have not yet heard a response, we collected practice SAT tests from a 2005 CollegeBoard preparation book. We evaluated on all reading comprehension sections from 3 full SAT tests. These books were already purchased several years ago. For convenience, we use "test" to refer to just the 3 reading sections of a test, and "reading section" to refer to just the reading comprehension passages and accompanying questions of each reading section.

Each practice test contains 3 reading sections. Within each section there can be *small* passages (about 100 words) with 4 questions, *medium* passages of 450-500 words and 5-8 accompanying questions, and *long* passages of 500-600 words with 12-13 questions. Of the first two sections, each has one small passage and 2 medium passages while the other has one small passage and 1 long passage. The third reading section always has one long passage. Each question has an associated difficulty as well (*easy*, *medium*, or *hard*). The Math and Writing sections, as well as the vocabulary questions in the reading section, are not used in evaluation.

Our practice tests are in paper format and needed to be converted to an electronic format. This conversion was actually quite complicated. First, every test was scanned into PDF format. Then these PDFs were converted into text format using a combination of Amazon Mechanical Turk and OCR software. For Mechanical Turk we split each test into individual passages - 2 small, 2 medium, and 2 large. Each passage is published as its own task for workers. The PDFs were converted into images and the images for each passage plus questions were uploaded to Amazon Web Services S3 File Storage. Each task consisted of the images on the left side of the screen and text boxes to type the text on the right side of the screen. With the OCR software, copies of each image had to be made, and the right-hand side and left-hand side of the image processed separately. Each SAT page has

two columns, which are misinterepreted by the OCR software as a single line. Tests 2 and 3 were processed via Mechanical Turk while Test 1 was processed via OCR software. The expenditure on Mechanical Turk was $116.92 and expenditure on OCR was $17.95, for a total of $134.87. Due to some errors in the conversion process, a few questions (mainly on small passages) were not used in the full evaluation. The passage text, question text, and answer key are all stored in text form and used as input for evaluation (Section 6.2).

6. Evaluation and Discussion

In this section we present and interpret our results on three full SAT reading comprehension tests. First we discuss the peformance of each model on each test and compare the performance of our enhanced models to their baseline versions. We break down what types of problems each model did well / poorly on, and offer an interpretation as to why. Then we compare the Tensor model to the Memory Networks model and evaluate their relative performance. Finally we propose future improvements.

6.1. Implementation Details

There are three main components to the implementation: knowledge graph creation, tensor decomposition, and neural network. We heavily modified the Java code provided by Narasimhan and Barzilay[1] [17], which itself uses the Stanford CoreNLP Toolkit[2] [15], to create the entity-relation triple extraction program. The RESCAL algorithm used is the version created by Nickel[3] [18]. We modify this code to include our customized regularization term. The memory network is implemented using Numpy[4], Theano[5], and Jupyter Notebook[6] [12, 20, 2, 4].

---

[1]https://bitbucket.org/ghostof2007/mctest
[2]http://stanfordnlp.github.io/CoreNLP/
[3]https://github.com/mnick/rescal.py
[4]http://www.numpy.org/
[5]http://www.deeplearning.net/software/theano/
[6]http://ipython.org/notebook.html

6.2. Evaluation Process

The evaluation process starts with the text input separated into passage text, questions, and answers. For each passage, the passage text is fed through the knowledge extraction program, which converts it into a set of entity-relation triples and outputs those into another text file.

The tests are answered passage-by-passage. This means that each model evaluates on one passage at a time. Each passage is considered its own knowledge graph. This means that in the tensor model a new tensor of varying dimensionality is created for each passage, and the accompanying questions can only be answered as queries on that particular graph. Combining the knowledge graphs for different sections would only serve to make it harder to answer the questions (and it is already hard enough!) Therefore it makes sense to run the tensor model on each passage separately. The same holds true for the Memory Networks model. The neural net has an additional training constraint. Due to the dimensionality constraints on the embedding matrices and the underlying architecture, it is not possible to combine the tensors from different passages.

In each evaluation we first train the model and then evaluate its performance on the questions for that passage. The Tensor model is trained using ALS (Section 4.2.4). The training process is independent of the actual performance of the model, i.e. there is no feedback, so we can test the model's performance on every question. The neural net is trained in a feedback loop using backpropagation as discussed in Section 2.2.1 and Section 2.2.2. In each passage we set aside $\frac{2}{3}$ of the questions for training and use the remaining $\frac{1}{3}$ as the test set. For this reason we do not evaluate the neural net on any of the Small passages.

We run the Tensor model 10 times with $d = 20$ and pick the model with best performance. We train the Neural Network model for 1000 iterations, evaluating its performance every 50 iterations, and pick the model with best performance. Ideally we would have a validation set separate from the test set, but because we have to train passage-by-passage the number of questions is already incredibly low for each model.

6.3. Results

Section 8.2 in the Appendix contains full evaluation data for each test. The tables contain each model's answer for every question, as well as the correct answer for that question and its associated difficulty. We evaluated Test 1 on our enhanced models as well as their baseline versions for comparison. Both models outperformed the baselines, so on Tests 2 and 3 we evaluated only on our enhanced models. The analysis of exactly how our models perform better is in Section 6.4.

On Test 1 our Tensor model with regularization (Section 4.2) scores 8 / 41 while the model without regularization scores 7 / 41. The 1-layer and 10-layer Memory Networks both scored 2 / 13. Hereafter "tensor model" refers to the tensor model with regularization, and "MemN2N" refers to the 10-layer neural network.

On Test 2 the tensor model scored 11 / 40 while the MemN2N scored 7 / 14.

On Test 3 the tensor model scored 9 / 41 while the MemN2N scored 3 / 14.

6.4. Comparison to Baseline

On Test 1 we compared our enhanced model to their respective baselines. We updated the tensor model with SSE regularization, so its baseline is the model without regularization. The baseline for the 10-layer MemN2N is the 1-layer MemN2N.

The tensor model with regularization scores 1 better than the model without regularization. Besides pure score, there is additional evidence that adding regularization creates a better model. For each passage we train the model 10 times and evaluate its performance. Without regularization, the training process produces the same model each time. However, with regularization it produces a slightly different model each time. The optimization process with ALS approximates the graph structure, which means that the tensor factorization should be slightly different each time. If it is the same, as is the case in the model without regularization, then the inherent properties are being found *too easily*, suggesting that it is not a good factorization of the graph.

The regularization term we use imposes some geometric constraints on the model. Training without regularization allows the training to violate these constraints, producing a model which does not actually represent the underlying graph. Adding the regularization term keeps the model flexible by forcing the factorization to more accurately represent the underlying graph. The model is not only more likely to answer correctly, but also more resistant to errors in the knowledge graph extraction. We discussed in Section 4.4 that the tensor model is too reliant on the knowledge graph extraction process. Preserving flexibility is of utmost importance, and the regularization term is necessary to do that.

With regards to the neural net, the 10-layer MemN2N is significantly better than the 1-layer version. Although both score the same on Test 1, the 1-layer MemN2N does not actually learn any of the underlying representation! Table 1 demonstrates that in each section, the 1-layer MemN2N predicted the same answer choice for every question. Its choice was the letter answer choice seen most in the training set. For example, in the Section 2 first Medium passage passage there are two questions where the answer is *C*. No other answer choice appears more than once. As a result, the 1-layer MemN2N guesses *C* for all of the remaining questions. The 10-layer MemN2N, on the other hand, guesses different answers for different questions in the same section. This indicates that it is predicting an answer based on some underlying semantic information, rather than simply the most commonly-seen answer choice. Simply guessing the most commonly-seen answer choice is not useful in evaluating how well the model understands the semantic information of the questions and the underlying knowledge graph. Therefore the 10-layer MemN2N is a much better model.

The reason the 1-layer MemN2N only outputs the most commonly-seen answer choice is that it overfits very quickly. Figure **??** shows how the 1-layer MemN2N rapidly fits the answer distribution to the training data, and becomes more confident in its answer with more training iterations. We can guard against overfitting by adding more layers.

6.5. Interpretation of Results

In this section we interpret the models' performance. We first note that the model performance is not consistent across passages, but rather very good on some passages and very poor on others. For example, on Test 2 Section 1 the tensor model scores 3 / 12 but on Test 2 Section 2 on the second Medium passage the tensor model scores 4 / 9. Similarly, on Test 2 Section 3 the MemN2N scores 4 / 5 but on Test 3 Section 2 the MemN2N scores 0 / 4.

There are several possible reasons for this phenomenon. The first is that both models are dependent on the knowledge graph extraction, so if the extraction process is accurate for a given passage then both models will perform well, and if the extraction is inaccurate then both models will perform poorly. This reason is unlikely, however, because the two models do not necessarily perform well or poorly in the same passages. On Test 2 Section 3 the MemN2N scores 4 / 5 but the tensor scores just 2 / 13. If the knowledge graph extraction were influencing the consistency of our models, they should both perform equally well or poorly in a given section.

The more likely reason, and the one that proves the usefulness of our models, is that within those passages the models are able to find the inherent properties of the knowledge graph (and therefore the text) and answer questions about it. Why the models work on some passages but not on others is unclear, but that is to be expected when working with machine learning algorithms. While the mathematics behind the tensor method are clear, the actual ALS process is still quite variable and hard to interpret. The neural network is even harder to understand, as its learning process is hidden behind 10 layers. However, we know that when the models are scoring very well on multiple sections it is not simply coincidence. That they score very poorly on multiple sections as well simply reinforces the idea that the models are answering questions based on their understanding of the text. Whether or not they can answer correctly depends on how well they parse the question and how accurate their representation of the text is. What dictates this accuracy is not easily seen due to the learned nature of the models. The goal of this research was to create models which would

be able to actually answer questions based on the semantic information, instead of just using bag-of-words matching or other basic techniques, and we have done just that.

This claim is strengthened when we investigate the passages which the models performed particularly poorly on. Test 1 Section 3 has perhaps the hardest passage for our models, with the tensor model scoring 0 / 13 and the MemN2N scoring 1 / 5. This passage is actually composed of two separate passages about the same topic, and most of the questions require comparing the opinions and viewpoints of the two authors. These types of comparison questions cannot be answered using the knowledge graph, because they require a higher level of analysis. Our models try to answer the question by completing a triple using the entities and relations on the knowledge graph, but the types of questions asked in Test 1 Section 3 require analyzing the entire knowledge graph itself. We expect our models to score poorly on such questions, and that is what happens. It is difficult to explain exactly what about certain passages makes our models score well, but we can understand which passages will cause our models to score poorly.

6.6. Performance by Difficulty

Figure 5 shows the total number of questions answered across all three tests sorted by their difficulty, while Figure 6 controls Figure 5 for the total number of questions of each difficulty.

The first observation we can make is that both models answer significantly more Medium questions than Easy or Hard, but that their percentage on Medium questions is lowest. This point is actually quite subtle. The full data tables demonstrate that there are far more Medium questions in general than Easy or Hard. However, the huge number of Medium questions means that the types of Medium questions are also quite varied. Whether the percentage of Medium questions answered correctly is so low is because there are simply far more Medium questions or because there are types of Medium questions that both models score poorly on (such as compare/contrast questions across passages) is unclear.

Figure 5: The total questions of each difficulty answered correctly by the models
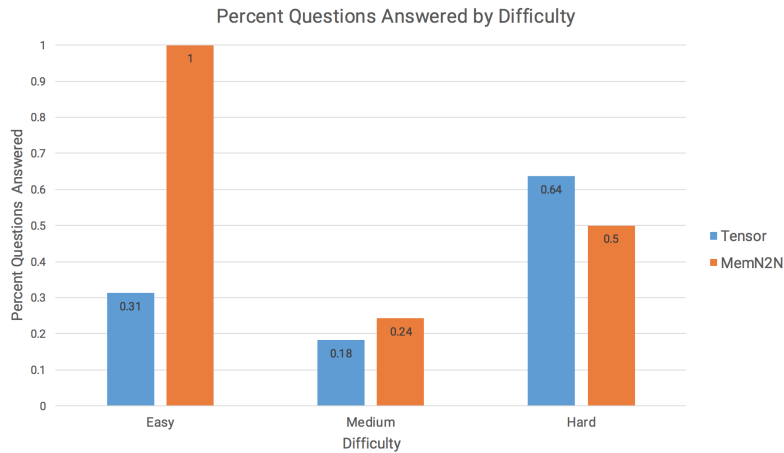
Figure 6: The percent of questions of each difficulty answered correctly by the models

The second observation we make is that the MemN2N model scores perfectly on the Easy questions and 50% on the Hard questions. The problem with analyzing this point is that the sample size is simply too small to make any claims about the MemN2N's performance on Easy or Hard questions. We can definitively say that the model was able to answer at least one of each type of question, even given the limited test size, but we cannot claim more than that.

The third important observation here is that the tensor model answers a very high percentage of Hard questions correctly. This is partially due to the relatively small sample size (only 11), but also a consequence of the way in which the tensor model is actually answering

these questions. Consider Test 3 Section 2, where the tensor model answers 2 out of 4 Hard questions correctly. This is actually another passage that contains two smaller passages, but both of the Hard questions the model gets right deal with only one of the passages. In fact, both questions ask about the meaning of a specific phrase in the context of the passage. The tensor model is actually able to match the words in the phrase with entities in the knowledge graph, and then match the answer choice with the relevant information about the matched entity.

As far as questions go, these are actually some of the easiest for the tensor model to answer because they do not require any high-level analysis. They simply require matching words in the question to entities in the graph, and then matching the corresponding graph element to words in the answer choice. This is exactly how the tensor model works, by using word2vec to match words in the question to entities in the graph, and then matching the relevant information to words in the answer choices.

This analysis leads us to a key observation. The difficulty standards for humans are not at all similar to the difficulty standards of our models. The Hard questions described in the paragraph above can be very difficult for human readers, because they require understanding the given phrase in context and then translating it into one of the answer choices. Essentially, these questions require fully understanding the phrase as well as its context. Our models, however, do not need to fully understand the phrase or its context; they simply need to match words in the question to entities in the graph. On the other hand, questions marked Easy can be very difficult for our models (namely the tensor model, as the MemN2N scored 2 / 2 on all its Easy questions). Some Easy questions ask about the author's possible opinion on some topic that is not directly contained in the text. For many humans these can be answered simply by understanding the author's tone throughout the passage. Our models cannot understand tone, and cannot process information not contained in the knowledge graph. These types of questions are therefore impossible for them to answer.

Figure 7: The percent of questions on each test answered correctly by the models

6.7. Comparing Tensor Model to MemN2N

Figure **??** shows the percent of questions answered correctly by each model on each test. There are several striking observations. The first is that the tensor model scores quite consistently, varying by just over 10% across all three tests. The MemN2N model is very volatile, in fact scoring the single lowest and highest performances across both models and all tests.

This result is to be expected, as we discussed in Section 4.4 how the tensor model is more mathematically sound while the neural net model does not have as clear mathematics describing how it works. The MemN2N's flexibility allowed it to score very high on Test 2, but at the same time score very low on Test 1.

We also note that the models scored relatively similar on all of the tests. They both scored lowest on Test 1, second-best on Test 3, and highest on Test 2. Additionally, on both Test 1 and 3, the scores for the tensor and MemN2N are almost identical, differing by no more than 2%. This result is partially due to the small sample size of MemN2N questions. Nevertheless, that the two models have similar scores across all three tests indicates that for a given knowledge graph they find the hidden representation of the knowledge graph equally well. We conclude that the two models are, for the most part, equally viable. Across

the three tests, the tensor model scores barely higher than the MemN2N model on Tests 1 and 3, but significantly worse on Test 2. The results demonstrate that the tensor model is the better general choice, as it guarantees some amount of competence, but at the same time it will not score very high. The MemN2N will most likely score slightly worse than the tensor model, but has much higher potential.

7. Future Work

We hope this work is just the start of research into answering SAT questions and tackling machine comprehension tasks of similar difficulty. To that end, we present several ways to improve the models discussed in this paper. These improvements fall in three major areas: linguistic analysis, tensor model, and neural network.

7.1. Linguistic Analysis

There are three instances in our model where we need to analyze the semantic or structural information of text. The first is extracting the knowledge graph triples from the passage text. The second is parsing the question and understanding what it is asking. The third is choosing an answer choice given the information output by the model - this applies only to the tensor model.

7.1.1. Knowledge Graph Extraction Improvements

There are some drawbacks to the way we currently extract the knowledge graph triples from the text. Some information is lost in the extraction process. Because the subject is forced to be $entity_1$ we lose information about the direct object. For example, in the sentence in Figure 3, we learn that the dog likes sausage but not that the sausage is spicy.

Additionally, SAT sentences are very difficult to parse. Most have multiple clauses, and normally unusual structural elements such as the colon or dash are quite commonplace. One sentence used in an actual passage is: *The critic Edumund Wilson was not a self-conscious letter writer or one who tried to sustain studied mannerisms.* We are able to extract the

triple ⟨Wilson, was, writer⟩ but due to the complicated structure of the sentence we are not able to extract that Wilson was not self-conscious or that Wilson did not *try to sustain studied mannerisms.* Overall we can extract the most important points, but we often miss the subtleties of the sentence.

If we can add more semantic rules to triple extraction then both the tensor model as well as the MemN2N could have more information at hand to answer the question. Right now we only focus on the *nsubj* and *dobj* or *cop* dependencies. One useful dependency we could add is the *amod* dependency, which is for adjectival modifiers [9]. This could help us capture descriptive information about entities that is not given with a copular verb. We can also focus on entities that are not the nominal subject, or add logic to handle sentences with multiple clauses.

7.1.2. Question Parsing

In our existing model, we vectorize a given question by using word2vec to calculate the similarity of every word in the sentence (excluding stopwords) to every possible entity and relation, and choosing the 2 entities or entity and relation that have highest similarity to question words. This works well when the question tests a specific concept, but not when the question is more complicated. This system runs into the same problem as bag-of-words matching, where the semantic and structural information of the sentence is lost.

A better way to tackle question parsing would be to feed the questions through the knowledge graph extraction program, and extract the information from the question. However, because we need to frame the question as a query on the knowledge graph we would modify the extraction program to extract triples with a missing entity or relation. We can then try to fill in these missing elements using the original knowledge graph, and match the resulting filled-in triples to the answer choices. The answer would be the choice with strongest match. Implementing this would require the improved deep extractor discussed above in Section 7.1.1.

7.1.3. Choosing an Answer

Currently in the tensor model we choose an answer similarly to the way we parse a question-by using word2vec to find the single word which matches the model's output best. Often, the answer choices are sentences too, so we need to be able to parse them. Just as with the questions we can feed the choices through the Knowledge Graph extractor to find a set of triples. We then need to create some metric to match triples and use it to pick the best choice.

7.2. Tensor Model Improvements

There are several improvements that can be made to the tensor model. The first is that the model's accuracy can be affected greatly by errors in the knowledge graph extraction process. We need to find a way to make the tensor model robust to such fluctuations. The best way to do this is to use an ensemble training method. We can take the knowledge graph tensor and create several variants, each slightly different from the original. Each of these are trained separately, and then put together to create the actual approximation. This approximation has several redundancies and is therefore less affected by errors in the original knowledge graph. Ensemble learning methods are common in many machine learning algorithms.

Another improvement that can be made to the tensor model is to provide supervision in the training process. Because the model is not trained on the actual questions, it cannot learn from the structure of the questions. By adding a small supervised training component, we can achieve the flexibility of the MemN2N while keeping the performance guarantee of the tensor model.

7.3. MemN2N Improvements

There are two main improvements we can make to the Memory Network. The first is with the underlying architecture itself. The input vectors must be embedded with matrices A, C. Since each passage has a tensor of different shape, we cannot reuse A, C across passages.

We discussed in Section **??** how it is advantageous to train on just one passage at a time. However, with the neural network there may be some underlying patterns across passages that we are not able to capture currently. If the dimensionality constraints on the architecture are loosened, the MemN2N model can be trained across passages and even across tests. This means that the model could be trained on one particular test, and run fully on another test.

The second possibility is to improve the model's consistency while preserving its flexibility. The biggest drawback of the MemN2N currently is that it occasionally has very poor performance. If we can train the model such that it will be able to answer certain types of questions with high probability then the model will be more reliable. For example, the tensor model is able to answer most questions that ask for the meaning of some phrase in context (Section 6.6). If the model is more reliable and consistent while still flexible, then it can become the best tool for machine comprehension. The very nature of neural network training makes it hard to create consistency, but if the MemN2N has the ability to train across tests then a large training corpus can help it achieve the desired consistency.

References

[1] B. Bader, R. a. Harshman, and T. G. Kolda, "Temporal Analysis of Semantic Graphs Using ASALSAN," *Seventh IEEE International Conference on Data Mining (ICDM 2007)*, pp. 33–42, 2007. [Online]. Available: http://ieeexplore.ieee.org/lpdocs/epic03/wrapper.htm?arnumber=4470227

[2] F. Bastien, P. Lamblin, R. Pascanu, J. Bergstra, I. J. Goodfellow, A. Bergeron, N. Bouchard, and Y. Bengio, "Theano: new features and speed improvements," Deep Learning and Unsupervised Feature Learning NIPS 2012 Workshop, 2012.

[3] J. Berant and P. Clark, "Modeling Biological Processes for Reading Comprehension," *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 1499–1510, 2014. [Online]. Available: http://allenai.org/content/publications/berant-srikumar-manning-emnlp14.pdf

[4] J. Bergstra, O. Breuleux, F. Bastien, P. Lamblin, R. Pascanu, G. Desjardins, J. Turian, D. Warde-Farley, and Y. Bengio, "Theano: a CPU and GPU math expression compiler," in *Proceedings of the Python for Scientific Computing Conference (SciPy)*, Jun. 2010, oral Presentation.

[5] C. M. Bishop, *Neural networks for pattern recognition*, 1995, vol. 92.

[6] A. Bordes, N. Usunier, A. Garcia-Duran, J. Weston, and O. Yakhnenko, "Translating embeddings for modeling multi-relational data," in *Advances in Neural Information Processing Systems*, 2013, pp. 2787–2795.

[7] K.-W. Chang, W.-t. Yih, B. Yang, and C. Meek, "Typed tensor decomposition of knowledge bases for relation extraction," *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 1568–1579, 2014.

[8] D. Chen and C. D. Manning, "A Fast and Accurate Dependency Parser using Neural Networks," *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, no. i, pp. 740–750, 2014. [Online]. Available: https://cs.stanford.edu/{~}danqi/papers/emnlp2014.pdf

[9] M.-C. De Marneffe, T. Dozat, N. Silveira, K. Haverinen, F. Ginter, J. Nivre, and C. D. Manning, "Universal stanford dependencies: A cross-linguistic typology." in *LREC*, vol. 14, 2014, pp. 4585–4592.

[10] S. Guo, Q. Wang, B. Wang, L. Wang, and L. Guo, "Semantically smooth knowledge graph embedding," in *Proceedings of ACL*, 2015, pp. 84–94.

[11] L. Hirschman, M. Light, E. Breck, and J. D. Burger, "{D}eep {R}ead: {A} Reading Comprehension System," *Proceedings of ACL*, pp. 325–332, 1999.

[12] E. Jones, T. Oliphant, P. Peterson *et al.*, "SciPy: Open source scientific tools for Python," 2001–, [Online; accessed 2016-04-23]. [Online]. Available: http://www.scipy.org/

[13] Y. Koren, R. Bell, and C. Volinsky, "Matrix Factorization Techniques for Recommender Systems," *Computer*, vol. 42, no. 8, pp. 42–49, 2009.

[14] A. Kumar, O. Irsoy, J. Su, J. Bradbury, R. English, B. Pierce, P. Ondruska, I. Gulrajani, and R. Socher, "Ask Me Anything: Dynamic Memory Networks for Natural Language Processing," *arXiv*, pp. 1–10, 2015.

[15] C. D. Manning, M. Surdeanu, J. Bauer, J. Finkel, S. J. Bethard, and D. McClosky, "The Stanford CoreNLP Natural Language Processing Toolkit," *Proceedings of 52nd Annual Meeting of the ACL: System Demonstrations*, pp. 55–60, 2014.

[16] T. Mikolov, K. Chen, G. Corrado, and J. Dean, "Distributed Representations of Words and Phrases and their Compositionality," *Nips*, pp. 1–9, 2013.

[17] K. Narasimhan and R. Barzilay, "Machine Comprehension with Discourse Relations," *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pp. 1253–1262, 2015. [Online]. Available: http://www.aclweb.org/anthology/P15-1121

[18] M. Nickel, V. Tresp, and H.-P. Kriegel, "A Three-Way Model for Collective Learning on Multi-Relational Data," *28th International Conference on Machine Learning*, pp. 809—-816, 2011.

[19] M. A. Nielsen, *Neural Networks and Deep Learning*. Determination Press, 2015. [Online]. Available: http://neuralnetworksanddeeplearning.com/index.html

[20] F. Perez and B. E. Granger, "Ipython: A system for interactive scientific computing," *Computing in Science Engineering*, vol. 9, no. 3, pp. 21–29, May 2007.

[21] M. Richardson, C. J. C. Burges, and E. Renshaw, "MCTest: A Challenge Dataset for the Open-Domain Machine Comprehension of Text," *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing (EMNLP 2013)*, no. October, pp. 193–203, 2013.

[22] S. T. Roweis and L. K. Saul, "Nonlinear dimensionality reduction by locally linear embedding," *Science*, vol. 290, no. 5500, pp. 2323–2326, 2000.

[23] S. Sukhbaatar, A. Szlam, J. Weston, and R. Fergus, "End-To-End Memory Networks," pp. 1–11, 2015. [Online]. Available: http://arxiv.org/abs/1503.08895

[24] Z. Wang, J. Zhang, J. Feng, and Z. Chen, "Knowledge graph embedding by translating on hyperplanes." in *AAAI*. Citeseer, 2014, pp. 1112–1119.

[25] J. Weston, A. Bordes, S. Chopra, T. Mikolov, and A. M. Rush, "Towards AI-Complete Question Answering: A Set of Prerequisite Toy Tasks," 2015. [Online]. Available: http://arxiv.org/abs/1502.05698

[26] J. Weston, S. Chopra, and A. Bordes, "Memory Networks," *International Conference on Learning Representations*, pp. 1–14, 2015. [Online]. Available: http://arxiv.org/abs/1410.3916

8. Appendix

8.1. Derivation of Output Layer Error

In this section we present the derivation for the output layer error of an RNN with cross-entropy loss and softmax at the output layer. Let us consider an output layer with n nodes. Define z_i as the quantity before activation function at the output layer, and h_i as the final output of node i.

$$h_i = \text{Softmax}(z_i) = \frac{e^{z_i}}{\sum\limits_{j=1}^{n} e^{z_j}} = \frac{e^{z_i}}{e^{z_i} + \sum\limits_{j \neq i} e^{z_j}} \tag{6}$$

The partial derivatives yield

$$
\begin{aligned}
\frac{\partial h_i}{\partial z_i} &= \frac{\sum\limits_{j=1}^{n} e^{z_j}(e^{z_i}) - e^{z_i}(e^{z_i})}{\left(\sum\limits_{j=1}^{n} e^{z_j}\right)^2} = \frac{e^{z_i} \sum\limits_{j=1}^{n} e^{z_j}}{\left(\sum\limits_{j=1}^{n} e^{z_j}\right)^2} - \left(\frac{e^{z_i}}{\sum\limits_{j=1}^{n} e^{z_j}}\right)^2 \\
&= h_i - h_i^2 = h_i(1 - h_i) \\
\frac{\partial h_i}{\partial z_j} &= \frac{\sum\limits_{j=1}^{n} e^{z_j} * 0 - e^{z_i}(e^{z_i})}{\left(\sum\limits_{j=1}^{n} e^{z_j}\right)^2} = \frac{-e^{z_i}}{\sum\limits_{j=1}^{n} e^{z_j}} \frac{e^{z_j}}{\sum\limits_{j=1}^{n} e^{z_j}} \\
&= -h_i h_j
\end{aligned} \tag{7}
$$

******WORK IN PROGRESS*******

8.2. Full Evaluation Data

Here we present the full evaluation data on three practice tests.

Table 1: Evaluation on Test 1

	Number	Tensor w/o Reg	Tensor w/ Reg	1-layer MemN2N	3-layer MemN2N	Correct Answer	Difficulty
Section 1							
Small	9	A	A	-	-	E	E
	10	A	D	-	-	B	M
	11	D	D	-	-	D	E
	12	E	E	-	-	D	E
Long	13	C	D	train	train	E	M
	14	D	D	train	train	D	E
	15	E	B	train	train	B	M
	16	E	A	train	train	A	H
	17	A	A	train	train	A	H
	18	E	D	train	train	C	E
	19	D	D	train	train	A	M
	20	D	A	train	train	B	M
	21	E	E	A	A	B	M
	22	A	A	A	A	C	M
	23	D	D	A	A	B	M
	24	A	B	A	A	A	E
Section 2							
Medium	13	B	B	train	train	C	M
	14	C	D	train	train	B	E
	15	E	E	train	train	C	M
	16	B	B	train	train	E	E
	17	A	A	C	C	D	E
	18	A	A	C	D	C	M
Medium	19	B	E	train	train	E	M
	20	E	E	train	train	C	E
	21	D	B	train	train	D	M
	22	E	E	train	train	E	M
	23	D	E	E	C	A	M
	24	B	B	E	C	B	M
Section 3							
Long	7	A	A	train	train	B	M
	8	C	C	train	train	A	M
	9	E	E	train	train	D	M
	10	D	D	train	train	E	M
	11	B	B	train	train	A	M
	12	C	C	train	train	A	E
	13	E	B	train	train	D	E
	14	B	B	train	train	D	M
	15	A	A	D	D	E	M
	16	A	A	D	D	E	M
	17	D	D	D	D	C	M
	18	D	D	D	A	A	M
	19	C	C	D	D	E	M
Total Correct		7	8	2	2		
Total Questions		41	41	13	13		

Table 2: Evaluation on Test 2

	Number	Tensor w/ Reg	3-layer MemN2N	Correct Answer	Difficulty
Section 1					
Long	13	C	train	A	M
	14	A	train	C	M
	15	D	train	E	M
	16	D	train	D	M
	17	B	train	D	M
	18	E	train	D	M
	19	E	train	E	M
	20	A	train	C	M
	21	B	E	A	M
	22	A	B	B	M
	23	B	B	B	H
	24	B	D	A	M
Section 2					
Medium	10	B	train	A	M
	11	B	train	A	M
	12	A	train	D	M
	13	C	train	C	E
	14	B	D	B	M
	15	E	A	A	M
Medium	16	A	train	E	M
	17	B	train	D	E
	18	A	train	A	E
	19	E	train	C	M
	20	B	train	E	H
	21	A	train	A	M
	22	C	E	C	M
	23	E	E	B	M
	24	B	E	B	M
Section 3					
Long	7	C	train	B	M
	8	E	train	B	E
	9	A	train	D	M
	10	C	train	D	M
	11	E	train	C	E
	12	E	train	A	M
	13	A	train	B	M
	14	E	train	E	H
	15	A	B	B	M
	16	E	E	E	M
	17	C	D	D	M
	18	A	B	B	M
	19	C	C	A	M
Total Correct		11	7		
Total Questions		40	14		

Table 3: Evaluation on Test 3

	Number	Tensor w/ Reg	3-layer MemN2N	Correct Answer	Difficulty
Section 1					
Medium	10	C	train	B	E
	11	E	train	C	M
	12	B	train	A	E
	13	A	train	D	M
	14	A	C	C	M
	15	C	C	D	E
Medium	16	A	train	A	M
	17	A	train	B	H
	18	B	train	E	M
	19	A	train	D	M
	20	B	train	D	M
	21	C	train	D	M
	22	A	E	A	M
	23	E	E	E	E
	24	C	E	B	M
Section 2					
Long	13	A	train	E	M
	14	E	train	C	M
	15	A	train	D	M
	16	D	train	D	H
	17	A	train	B	M
	18	C	train	C	M
	19	B	train	B	H
	20	E	train	B	M
	21	B	B	D	H
	22	A	B	E	H
	23	A	E	A	M
	24	D	A	E	M
Section 3					
Long	7	B	train	D	E
	8	E	train	E	M
	9	B	train	D	M
	10	A	train	D	M
	11	A	train	B	M
	12	E	train	C	E
	13	A	train	A	M
	14	D	train	E	M
	15	A	E	C	M
	16	D	B	B	M
	17	C	D	E	M
	18	B	D	A	M
	19	E	D	D	M
Total Correct		9	3		
Total Questions		41	14		