

Solving SAT Reading Comprehension Questions with Memory Networks

Saahil Madge

Adviser: Professor Christiane Fellbaum

1. Background and Motivation

Machine question-answering (QA) has been a popular subject in recent research. It is a very interesting field from an academic perspective, requiring knowledge and techniques across the areas of artificial intelligence, machine learning, natural language processing, linguistics, and mathematics.

QA has applications to many aspects of daily life. A famous use of question-answering techniques is in Apple's Siri program. Siri combines speech recognition software with question-answering techniques to create a virtual iPhone assistant. Another example is IBM's Watson, which answered questions in many subjects and performed quite well as a contestant on the game show *Jeopardy!* A less spectacular but more commonly-used example can be seen by typing in a simple question such as *How many calories are in an apple?* on Google. The search results provide a detailed nutritional breakdown of an apple.

The general term QA covers a broad array of subfields. Information retrieval, factual question-answering, reading comprehension, and inference are all part of the greater question-answering area. QA is a rapidly evolving field, but as the examples above indicate, recent research has focused extensively on answering spoken questions, and factual questions.

However, the subfield of machine comprehension has not been studied nearly as well, though it has become quite popular in the past few years. Machine comprehension involves training programs and systems to read text and answer questions based on the newly acquired information. In most other QA tasks, the system has a large database of facts (knowledge base) and must interpret the question and provide the relevant factual answer. Machine comprehension tasks, on the other hand, require the system to parse both the informational text as well as the question, and then provide the correct answer. As there is possibility for error within the knowledge base, machine comprehension tasks are some of the most difficult QA tasks.

Machine comprehension has a wide range of applications. Reading comprehension tests are a natural choice, but fields like document retrieval also use machine comprehension techniques. Financial firms that use news analysis to make trading decisions heavily rely on these techniques as well. Many long-term goals of AI such as dialogue and humanoid robots cannot be achieved without significant advance in this area [14]. Another incentive to study machine comprehension is that new research here can be applied to drive research in many other areas of artificial intelligence and natural language processing.

As a consequence of machine comprehension tasks being so difficult, most research with reading comprehension tests has focused on elementary school level tests. The most common dataset is MCTest[12], which we discuss in more detail in 3. This dataset is a set of passages and reading

comprehension questions created by crowdsourcing. The stories are “carefully limited to those a young child would understand”[12]. Another common dataset is Facebook’s bAbi dataset[14], which has a collection of short toy problems that can be used to train machine comprehension systems. These problems could certainly be solved by a human elementary school student. Berant et al.[2] train a system to answer questions based on a paragraph describing biological processes, which is certainly an advanced topic, but the system is specialized for biological tasks and not applicable as is to general reading comprehension.

In this paper we focus on SAT reading questions. The major reason is that they are far more complex than MCTest questions and other commonly-used machine comprehension datasets. The sentence structures, vocabulary, and topics covered in SAT passages are all far more varied than in the other benchmarks. The questions are much harder as well, and simply understanding the question and the individual choices is difficult by itself. SAT questions often test underlying themes and broad ideas rather than particular factual details (“Why” or “How” questions). In contrast, MCTest questions are often of a “Who”, “When”, or “What” form, which have traditionally been easier for programs to answer.

Another aspect that makes SAT questions so difficult is that they require inference to answer, as opposed to just matching words or syntax. MCTest and similar questions can mostly be answered through word-matching (discussed in 2.1.1) or via syntactic matching. Word-matching means the system will try to find which sentences in the passage have the most words in common with the query, and use that to choose an answer. Syntactic matching is the same principle, but aims to match syntactic structure. SAT questions, however, almost never repeat the answer words in the question itself, and the query structure is not correlated with passage structure. To solve these questions, our system has to really understand the high-level concepts and information that is implied but not directly on the surface. This is a significantly harder problem than has been tackled before, but it is necessary to solve to continue moving forward in the field of machine comprehension. Essentially, by trying to answer SAT questions we can try to solve problems that will be present in real-world applications but have not yet been tackled by existing research.

In particular, we combine two techniques: Memory Networks and relational semantic analysis. Memory Networks were introduced by Weston et al. in [15] and improved by Sukhbaatar et al. in [13]. These are a particular type of neural network. We discuss memory networks and alternatives in much greater depth in 3 and 6. Most recent machine comprehension research has focused on using neural nets as the core system, so we use them in this paper as well. Neural nets are a very generalizable technique and in the past decade have become very popular. In fact, advances in neural net techniques have greatly contributed to the recent surge in machine comprehension research.

Memory Networks require the input text and the query to be embedded in some form for easy computation. These embedding techniques are covered in detail in ???. Typically the embedding is done with bag-of-words embedding. This embedding technique relies on the same word being used in the answer choices as well as in the original text, which is true for MCTest and tests of similar difficulty. However, SAT reading comprehension passages often have very complex answer choices that rarely use the exact same word as used in the text. As a result we have to change the way we represent and parse the text. Instead of using *one-hot* vector encoding and picking answer choices based on how many words are in common with the query or text, we represent the text using relational logic. Essentially we dynamically create a “Knowledge Graph” as we parse through the text and try to frame each question as a query on this knowledge graph. We enhance a tensor factorization method known as RESCAL [1][10] and use this enhanced method to embed the text,

questions, and answers. Once the embedding is done we can train our Memory Network on the text and query to find the answer. The knowledge graph representation is typically used for factual data that has already been collected in entity-relation triple form. As far as we know, no previous research in this space has tried to dynamically create a knowledge graph on which we can answer questions.

Our goal with this research is to present a way to combine ideas that were previously used in isolation for QA or other learning tasks into a system which can parse and interpret a complex text, and answer difficult high-level questions about the text. The text, questions, and techniques we use are more ambitious than those studied in previous research. Our results should be interpreted as merely proof-of-concept. We hope that this paper will encourage others to tackle problems of equal or greater ambition, and drive progress in the development of safe artificial intelligence.

2. Technical Review

In this section we review some of the techniques that are used in previous research and in this paper.

2.1. Embedding Techniques

We begin by reviewing various methods for embedding text and performing queries on the embedded representations.

2.1.1. Bag-of-words Bag-of-words is a type of analysis in which we treat each sentence as simply a collection of words, without paying attention to their order or the dependencies between individual words. If we have some vocabulary of size V , we can encode each sentence as a vector of size V where each vector element corresponds to some word in the vocabulary. If the word appears in the sentence we store how many times it appears, and if it does not appear we store it as a 0. This is known as *one-hot* vector encoding. Most designs encode each input sentence and the query as one-hot vectors. These vectors are not used as-is, but are first embedded into some dimension d . The learning algorithm is run on the embedded vectors, and we can output a response in a specified format. Two common methods are to return a probability vector of size V where each element represents how likely that word is to be the output, or to pick the most likely word from that distribution and return only that word.

2.1.2. word2vec The *word2vec* embedding technique introduced by Mikolov et al. [8] is perhaps the most common embedding technique used in NLP research. It is actually an umbrella term for various methods, such as *skip-gram* or even bag-of-words embedding. The *word2vec* technique produces a vector representation of words, often with several hundred elements. These representations can then be manipulated as vectors in the higher-dimensional space, allowing for interesting operations on words. For example, $\text{vector}(\text{"king"}) - \text{vector}(\text{"man"}) + \text{vector}(\text{"woman"})$ produces a vector that is very close to $\text{vector}(\text{"queen"})$. The reason *word2vec* is able to embed so successfully is generally thought to be the fact that the process of embedding gives similar words a similar embedding (which is measured by cosine similarity).

2.1.3. Knowledge Graph A knowledge graph generally refers to a set of *entity-relation triples*. These triples take the form $\langle e_1, r, e_2 \rangle$, where e_1 and e_2 are entities and r describes the relation between these two entities. Entities and relations can be repeated across triples. The total set of

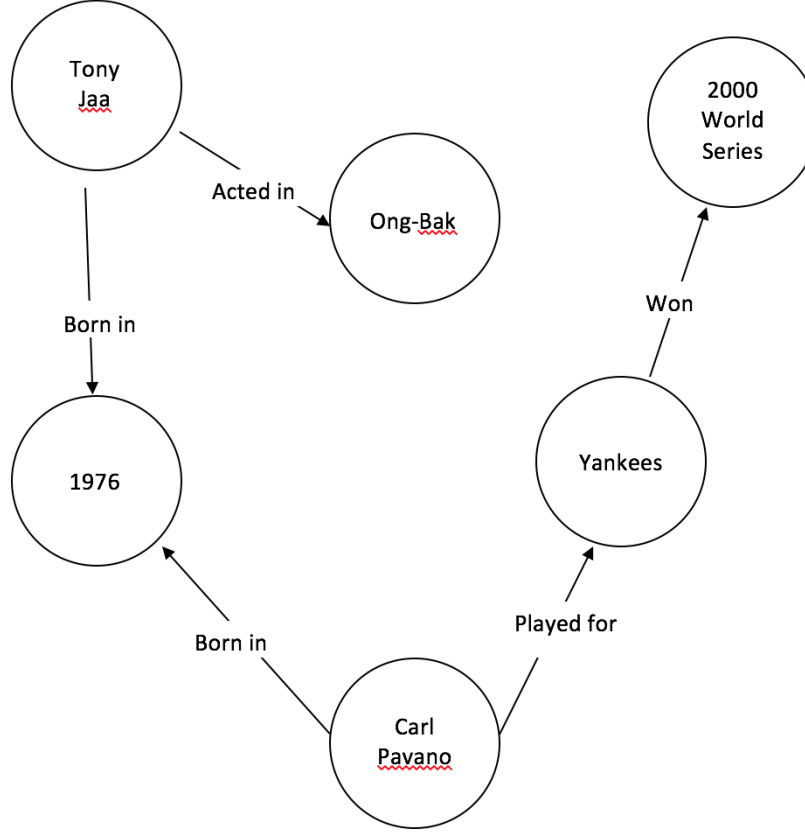


Figure 1: A small example knowledge graph

triples denotes all the knowledge in our “Knowledge Base”.

We can represent these triples as a directed graph, with each entity as a vertex and the relations as edge types between them. Figure 1 shows an example of a very small knowledge graph. Note that there is only one node per entity. The mathematical details of our knowledge graph model can be found in 6.

2.2. Neural Networks

As discussed in 1, recent research in machine comprehension has extensively used neural nets (NNs) and their variants. Here we provide a quick review of neural network design. These topics can be covered in far more depth in [3] and [11].

The most basic model is known as a *feed-forward* neural net. These have several layers composed of units (nodes). There is an initial input layer, a final output layer, and *hidden layers* in the middle. A unit in layer l takes as input the output of nodes in the previous layer (or the actual input if this is the input layer). It multiplies these by some weight matrix W^l , transforms it using an activation function g , and sends its output to the next layer (or as the final output if this is the output layer). Common activation functions are the *tanh* and *sigmoid* functions. The sigmoid function is defined as $\sigma(x) = \frac{1}{1+e^{-x}}$. For now we will assume that our activation function is sigmoid.

Mathematically, we define the input to the current layer as x^{l-1} , where x_{ij}^{l-1} means that it is the input from unit j in the previous layer to unit i in the current layer. W_{ij}^l is the weight on this connection, defined as 0 if node i does not depend on node j . The output of node j is

$\sigma(W^l x_j^{l-1})$. Typically the quantity $W x_j^{l-1}$ is known as z_j . We can also define bias vectors b such that $z_j = W^l x_j^{l-1} + b_j^l$. To summarize, we can write the output of some layer l in vector notation as

$$o^l = \sigma(W^l x^{l-1} + b^l)$$

We can propagate this forward starting at the input layer, through the hidden layers, and finally at our output layer. We can then train our network using the backpropagation algorithm. We have some cost function C that is a measure of the error of our output. We find the error at the output layer, and then change the weights at each previous layer, working backwards to the input layer. The key to backpropagation relies on updating the weights based on their partial derivatives with relation to C . In-depth derivations are provided in [3] and [11].

2.2.1. Recurrent Neural Networks Feed-forward networks are good tools, but do not have any form of state (memory). Here we review recurrent neural networks (RNNs) which have a hidden state h_t .

At step t in a recurrent neural network, we have input x_t and a hidden state h_{t-1} . We also have weight matrices U and W . We can define the current state $h_t = \sigma(Ux_t + Wh_{t-1})$. We have another weight matrix V for the output. The output $y_t = \text{softmax}(h_t)$, where $\text{softmax}(z_i) = \frac{e^{z_i}}{\sum_j e^{z_j}}$. The softmax function output is a probability distribution over the vector.

Using the softmax function has several advantages. Because it is a smooth function, we can take its derivative, making it easier to calculate the backpropagation equations. Additionally, most RNNs and variants use a loss function known as cross-entropy loss. If the model output is \hat{y} and the real (target) is y , then cross-entropy loss is defined as $\sum_i y_i \log(\hat{y}_i)$. When we pair a softmax output with cross-entropy loss, the error at the output node i is simply $\hat{y}_i - y_i$. This is very convenient because it can be calculated rapidly, so it allows for fast training of the RNN. Training is normally a bottleneck, and by combining a robust output function and robust loss function we get a very elegant output error. As this is a very important result, we have included it in 7.1.

3. Related Work

The first major research in machine comprehension was conducted by Hirschman, et al. [6]. Their system “Deep Read” takes a story and a series of questions as input, and proposes answers by using bag-of-words techniques with other methods such as stemming and pronoun resolution layered on top. On a collection of remedial reading questions for grades 3-6, Deep Read answered about 40% of the questions correctly.

Most recent research has used the MCTest [12], released by Richardson, et al. at Microsoft research. There are a total of 500 passages, each with 400 questions, created at the reading level of a young child. The stories are fictitious, which means the answers can only be found in the story itself. Richardson et al. also provide two baseline implementations to serve as a benchmark. The first is a simple bag-of-words model with a sliding window. It scores about 51% accuracy on the MC500 questions. The second implementation adds a distance-based scoring metric and reaches 56% accuracy. Both implementations score significantly higher on questions where the answer is contained in just one sentence than on questions where the answer requires information across multiple sentences.

Narasimhan and Barzilay [9] also work on the MCTest tasks. Their main insight is to create a task-specific discourse parser to capture specific discourse relations in the MCTest passages, while

the prevailing method had been to use generalized off-the-shelf discourse parsers. They create three models of increasing complexity. The first assumes each question can be answered using just one sentence, and estimates a joint probability distribution across the question, query, and answer. The second model adds in the joint distribution across a second sentence as well, to handle the case in which the answer needs two sentences to answer. The third model incorporates discourse relations between the two sentences to better understand the relationship between information in each sentence. The relations are defined as “Causal”, “Temporal”, “Explanation”, and “Other”. Most systems have performed the worst on these types of questions. By focusing specifically on modeling these explanatory relations, the third model easily outscores the other two, and the two MCTest baseline implementations, achieving almost 64% accuracy.

Berant et al.[2] also focus on analyzing inter-sentence relations, with application specifically to passages and questions about biological processes. These passages generally describe a chemical reaction or other process in which there are various starting entities which interact with each other and form a new output. Understanding how the inputs interact and tracing the flow of the process is crucial to answering the question. To solve this, Berant et al. define events (or non-events) as “triggers”, and try to find relationships between events. The events can be thought of as nodes on a graph, with an edge defining some relation between the two. There are eight possible relations, including “cause”, “enable”, and “prevent”. They first create events and then predict relations between the events. The queries are also formulated as a graph. They are categorized as dependency questions, temporal questions, or true-false questions. This model scores almost 67% on the dataset, over 6% better than the next-best model.

Weston et al.[15] introduced memory networks. These are a type of neural network which simulate a long-term memory. We discussed in 2.2.1 how RNNs have a “vanishing gradient” problem, and are not able to take advantage of states from more than a few steps prior. Memory networks have a memory module which stores old memories and then updates them given new inputs. When given a query, the memory network finds relevant memories, then finds memories that are relevant given those memories, and so on. Finally, it provides an answer to the query. Sukhbaatar et al.[13] improved this model by creating a model that can be trained end to end, while in the original design each module needed to be trained independently. This is the model we use for our design, so we discuss it in more detail in 6. Kumar et al.[7] create “Dynamic Memory Networks”. Their design uses two memory modules. The semantic memory module stores general knowledge, while the episodic memory module iteratively finds memories relevant to the query.

4. Data

For preliminary testing, we use both the MCTest [12] dataset as well as the Facebook bAbi dataset [14]. The MCTest baseline algorithms provide a benchmark to test our preliminary bag-of-words implementations against. The bAbi dataset is used to evaluate [13]. We use the same model so our preliminary neural net implementation is also evaluated on the bAbi dataset. However, we do not use as many training optimizations. These comparisons are meant to be benchmark comparisons, rather than trying to perform better.

The crux of our project relies on SAT reading comprehension tests for both training and testing. As there are no publicly available datasets, we tried contacting ETS to obtain a research dataset. As we have not yet heard a response, we also collected practice SAT tests from prep books. As of now we have 8 tests from a CollegeBoard prep book, and 11 from a Princeton Review prep book (10 full

practice tests and 1 practice PSAT test). We already owned these prep books.

Each practice test contains 3 reading sections. Each reading section has approximately 18 comprehension questions. About 4-6 of these are from short passages (100 words), and the remaining are from longer passages (450-500 words). There are 2 short passages and 1-2 long passages per section, each followed by comprehension questions. Hereafter we use “test” to refer to just the 3 reading sections of a test, and “reading section” to refer to just the reading comprehension passages and accompanying questions of each reading section. Math and writing sections, as well as the vocabulary questions in the reading section, are ignored.

Our practice tests are in paper format, so we must put them in an electronic format. Amazon Mechanical Turk was used for this task. Each test was scanned, and we approximated that typing up one test requires an hour. We requested Master Turkers, 2 per test, and paid \$8.00 for each task. The total expenditure was \$304.00.

5. Progress

Besides all of the preliminary research, the basic neural network model has been implemented, in both 1-hop and 3-hop version, as specified in [13]. Additionally, the programs to read in and format data from MCTest, bAbi, and the SAT tests have been written and tested. The tests have been scanned and the Turk tasks should be finished within a week.

The remaining steps are to create a model for dependencies within sentences and across sentences. Two approaches look promising. The first is to use a joint probability model as defined in [9] and with dependency parsing by [5]. The second is to try and create a concept map that is updated as we parse the passage. The questions can then be answered by looking at the concept map for a particular entity. This is similar to [4]. Once we have that we can combine the neural net with the parsing to train our final model.

6. Model

References

- [1] B. Bader, R. a. Harshman, and T. G. Kolda, “Temporal Analysis of Semantic Graphs Using ASALSAN,” *Seventh IEEE International Conference on Data Mining (ICDM 2007)*, pp. 33–42, 2007. [Online]. Available: <http://ieeexplore.ieee.org/lpdocs/epic03/wrapper.htm?arnumber=4470227>
- [2] J. Berant and P. Clark, “Modeling Biological Processes for Reading Comprehension,” *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 1499–1510, 2014. [Online]. Available: <http://allenai.org/content/publications/berant-srikumar-manning-emnlp14.pdf>
- [3] C. M. Bishop, *Neural networks for pattern recognition*, 1995, vol. 92.
- [4] A. Bordes, N. Usunier, S. Chopra, and J. Weston, “Large-scale Simple Question Answering with Memory Networks,” 2015. [Online]. Available: <http://arxiv.org/abs/1506.02075>
- [5] D. Chen and C. D. Manning, “A Fast and Accurate Dependency Parser using Neural Networks,” *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, no. i, pp. 740–750, 2014. [Online]. Available: <https://cs.stanford.edu/{~}danqi/papers/emnlp2014.pdf>
- [6] L. Hirschman, M. Light, E. Breck, and J. D. Burger, “{D}EEP {R}EAD: {A} Reading Comprehension System,” *Proceedings of ACL*, pp. 325–332, 1999.
- [7] A. Kumar, O. Irsoy, J. Su, J. Bradbury, R. English, B. Pierce, P. Ondruska, I. Gulrajani, and R. Socher, “Ask Me Anything: Dynamic Memory Networks for Natural Language Processing,” *arXiv*, pp. 1–10, 2015.
- [8] T. Mikolov, K. Chen, G. Corrado, and J. Dean, “Distributed Representations of Words and Phrases and their Compositionality,” *Nips*, pp. 1–9, 2013.
- [9] K. Narasimhan and R. Barzilay, “Machine Comprehension with Discourse Relations,” *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pp. 1253–1262, 2015. [Online]. Available: <http://www.aclweb.org/anthology/P15-1121>

- [10] M. Nickel, V. Tresp, and H.-P. Kriegel, “A Three-Way Model for Collective Learning on Multi-Relational Data,” *28th International Conference on Machine Learning*, pp. 809—816, 2011.
- [11] M. A. Nielsen, *Neural Networks and Deep Learning*. Determination Press, 2015. [Online]. Available: <http://neuralnetworksanddeeplearning.com/index.html>
- [12] M. Richardson, C. J. C. Burges, and E. Renshaw, “MCTest: A Challenge Dataset for the Open-Domain Machine Comprehension of Text,” *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing (EMNLP 2013)*, no. October, pp. 193–203, 2013.
- [13] S. Sukhbaatar, A. Szlam, J. Weston, and R. Fergus, “End-To-End Memory Networks,” pp. 1–11, 2015. [Online]. Available: <http://arxiv.org/abs/1503.08895>
- [14] J. Weston, A. Bordes, S. Chopra, T. Mikolov, and A. M. Rush, “Towards AI-Complete Question Answering: A Set of Prerequisite Toy Tasks,” 2015. [Online]. Available: <http://arxiv.org/abs/1502.05698>
- [15] J. Weston, S. Chopra, and A. Bordes, “Memory Networks,” *International Conference on Learning Representations*, pp. 1–14, 2015. [Online]. Available: <http://arxiv.org/abs/1410.3916>

7. Appendix

7.1. Derivation of Output Layer Error